

PARALLEL PROGRAMMING

An Introduction by Nixon



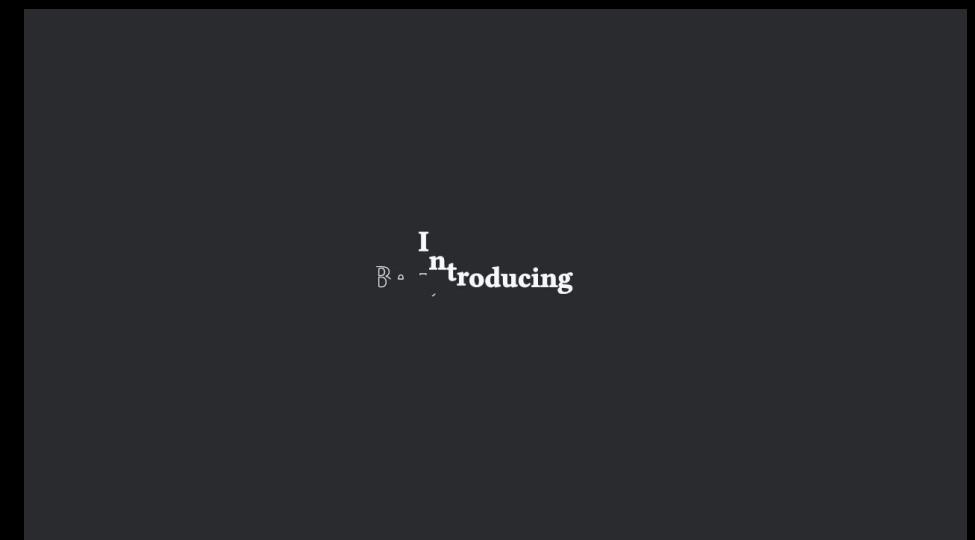
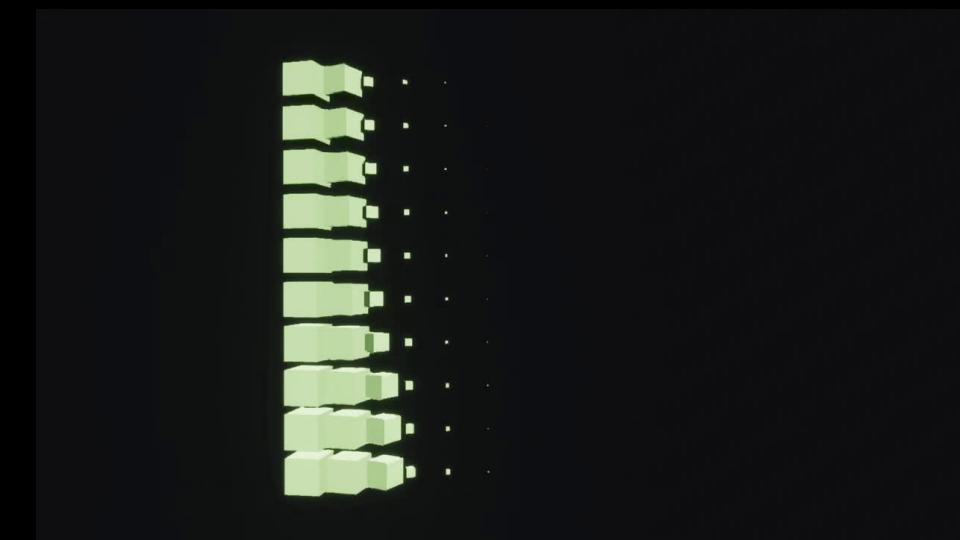
VOXELL





About Me

- Computer Graphics Engineer
- AI Enthusiast



VOXELL



Let's start off with a
question...

If you were plowing a field, which would you rather use?

2 strong oxen



1024 chickens

or



If you were plowing a field, which would you rather use?

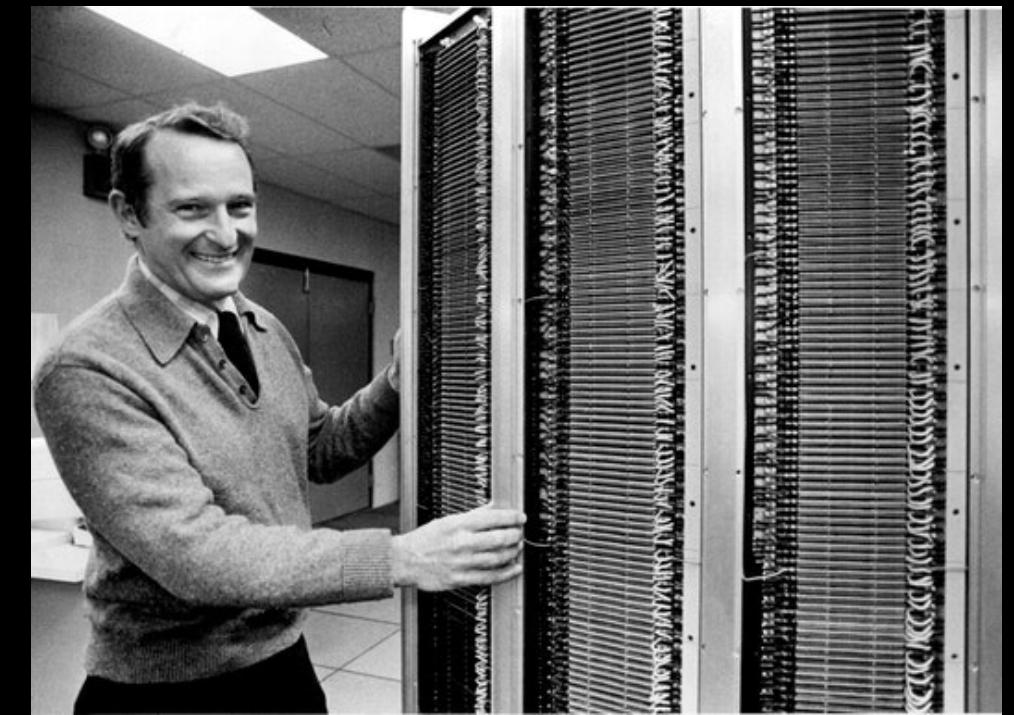
2 strong oxen



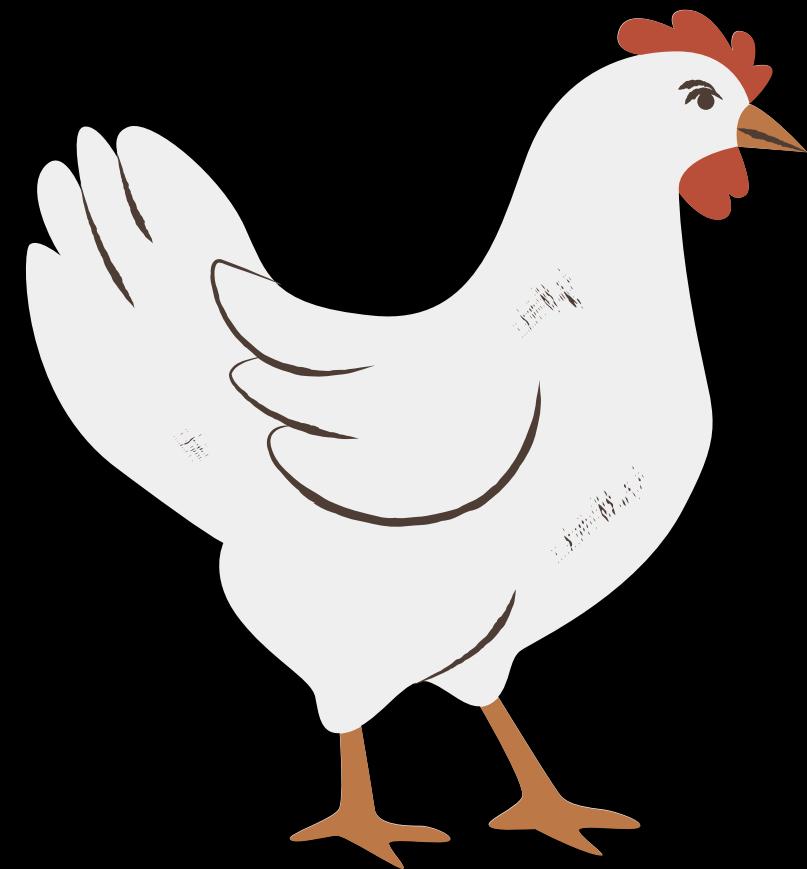
1024 chickens



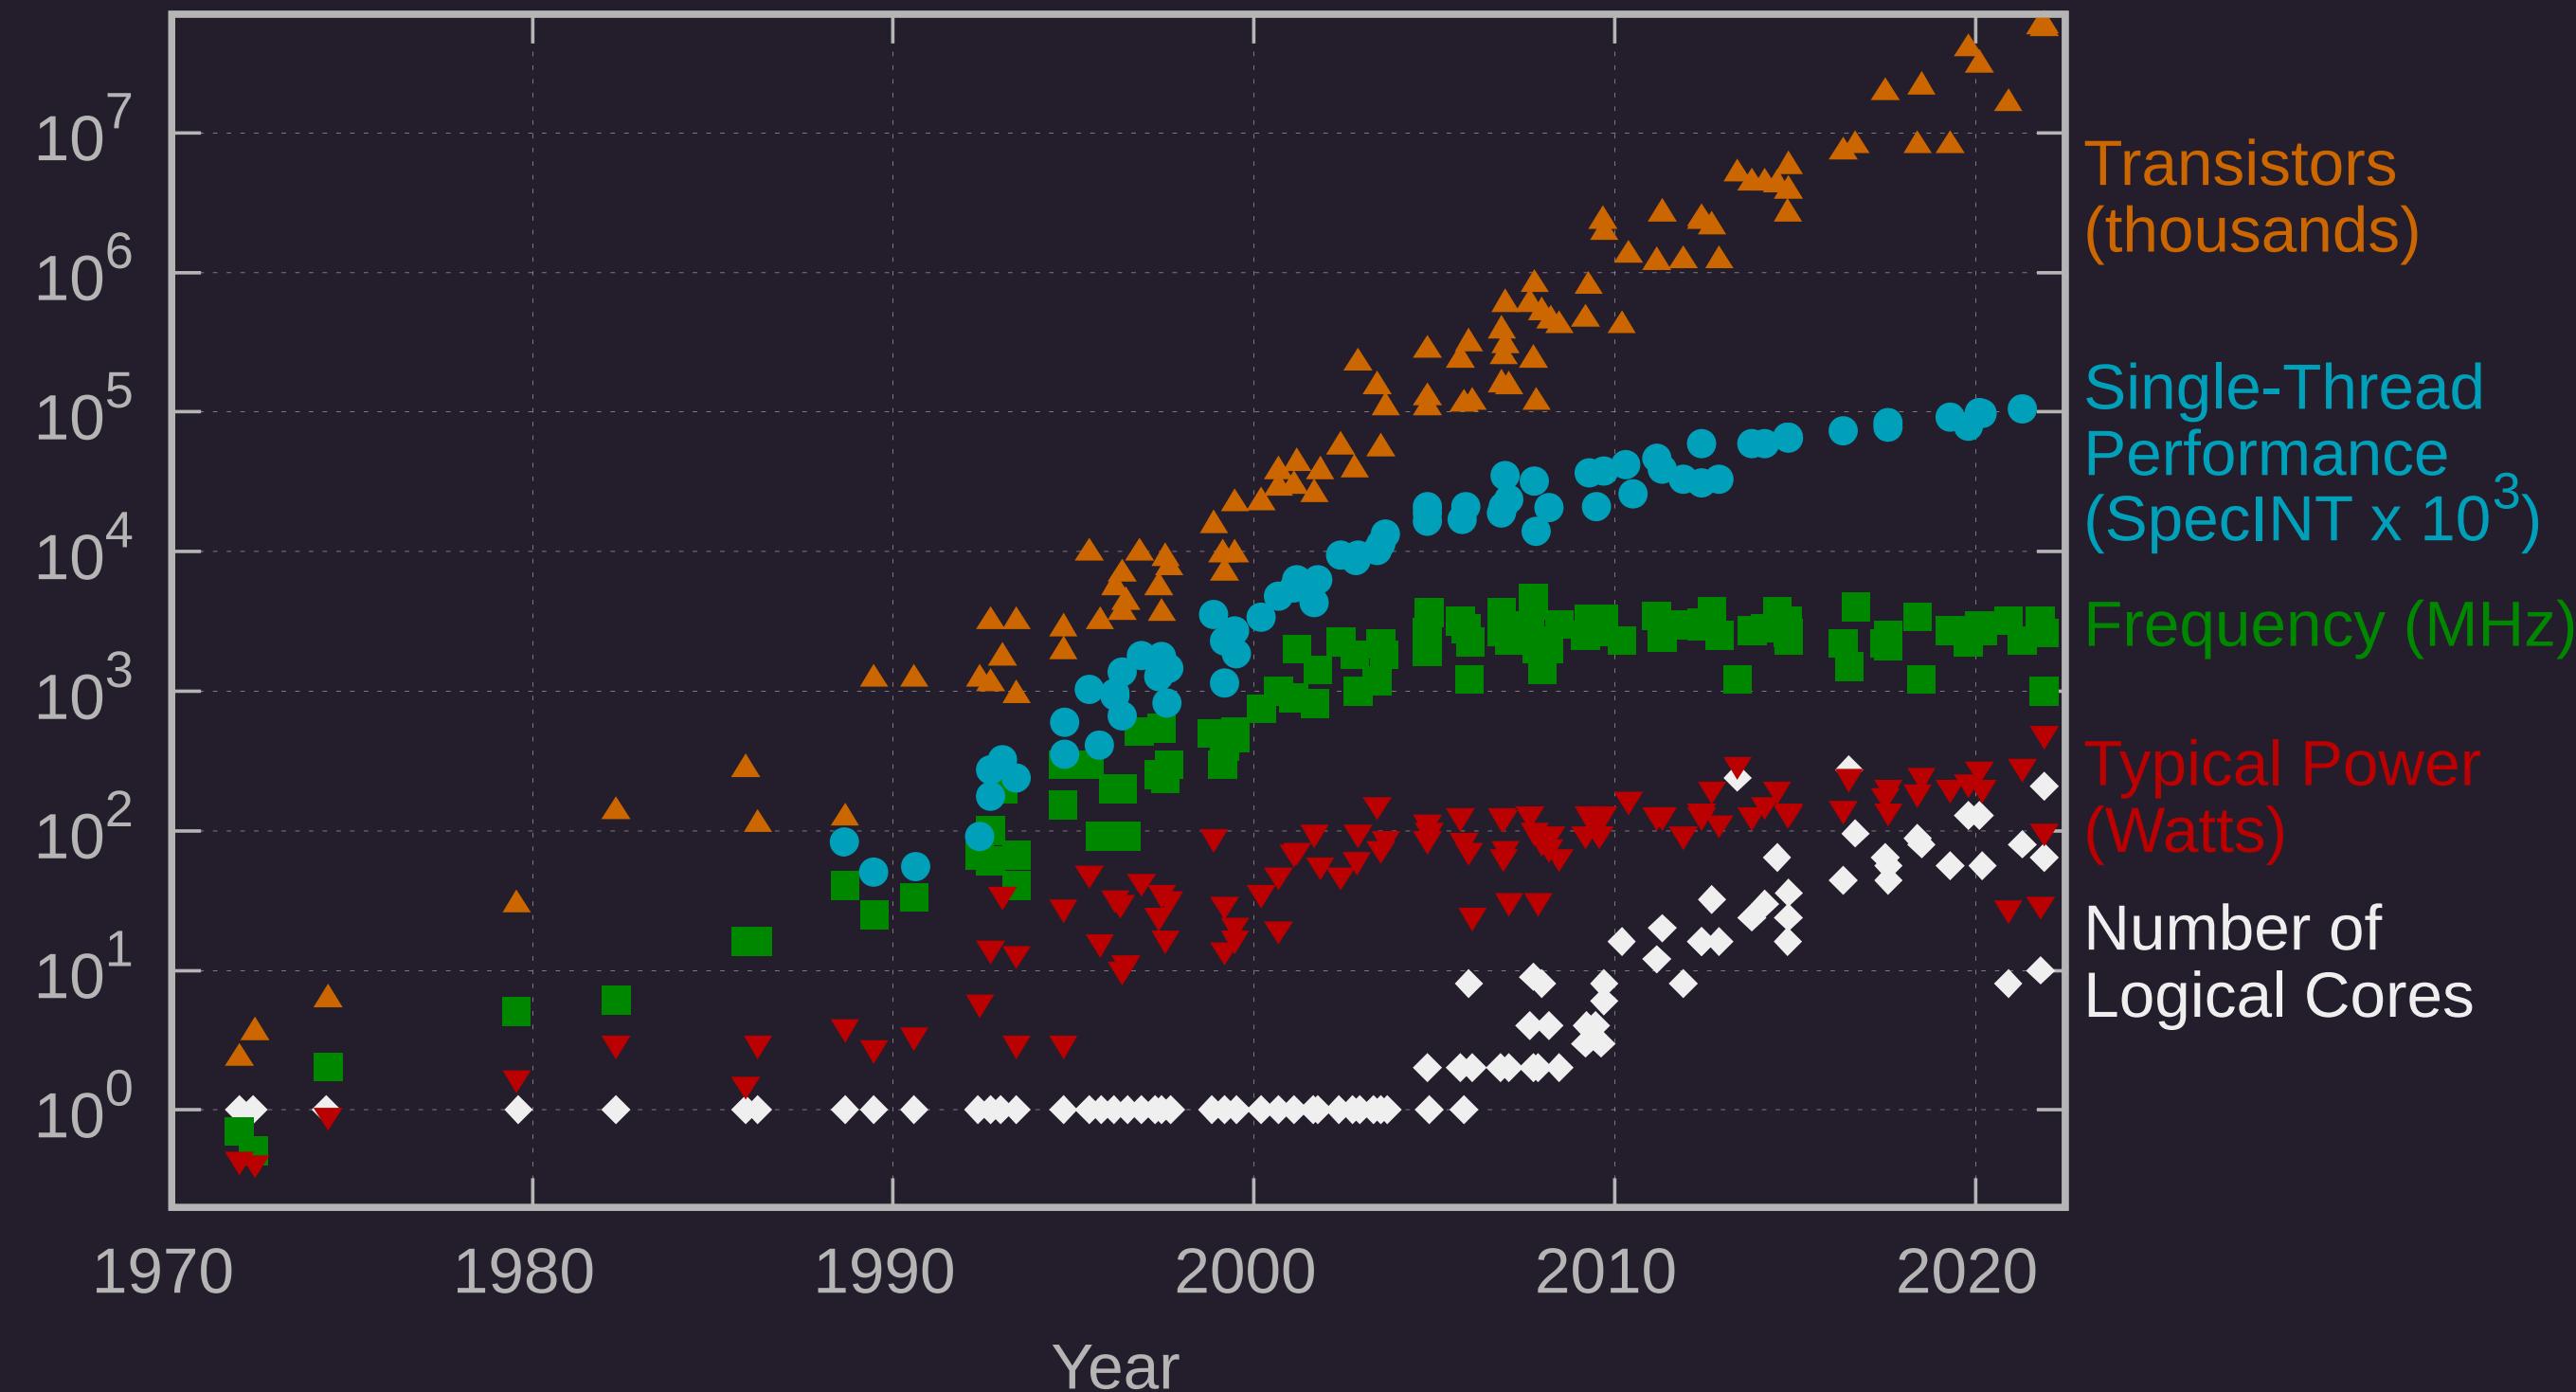
or



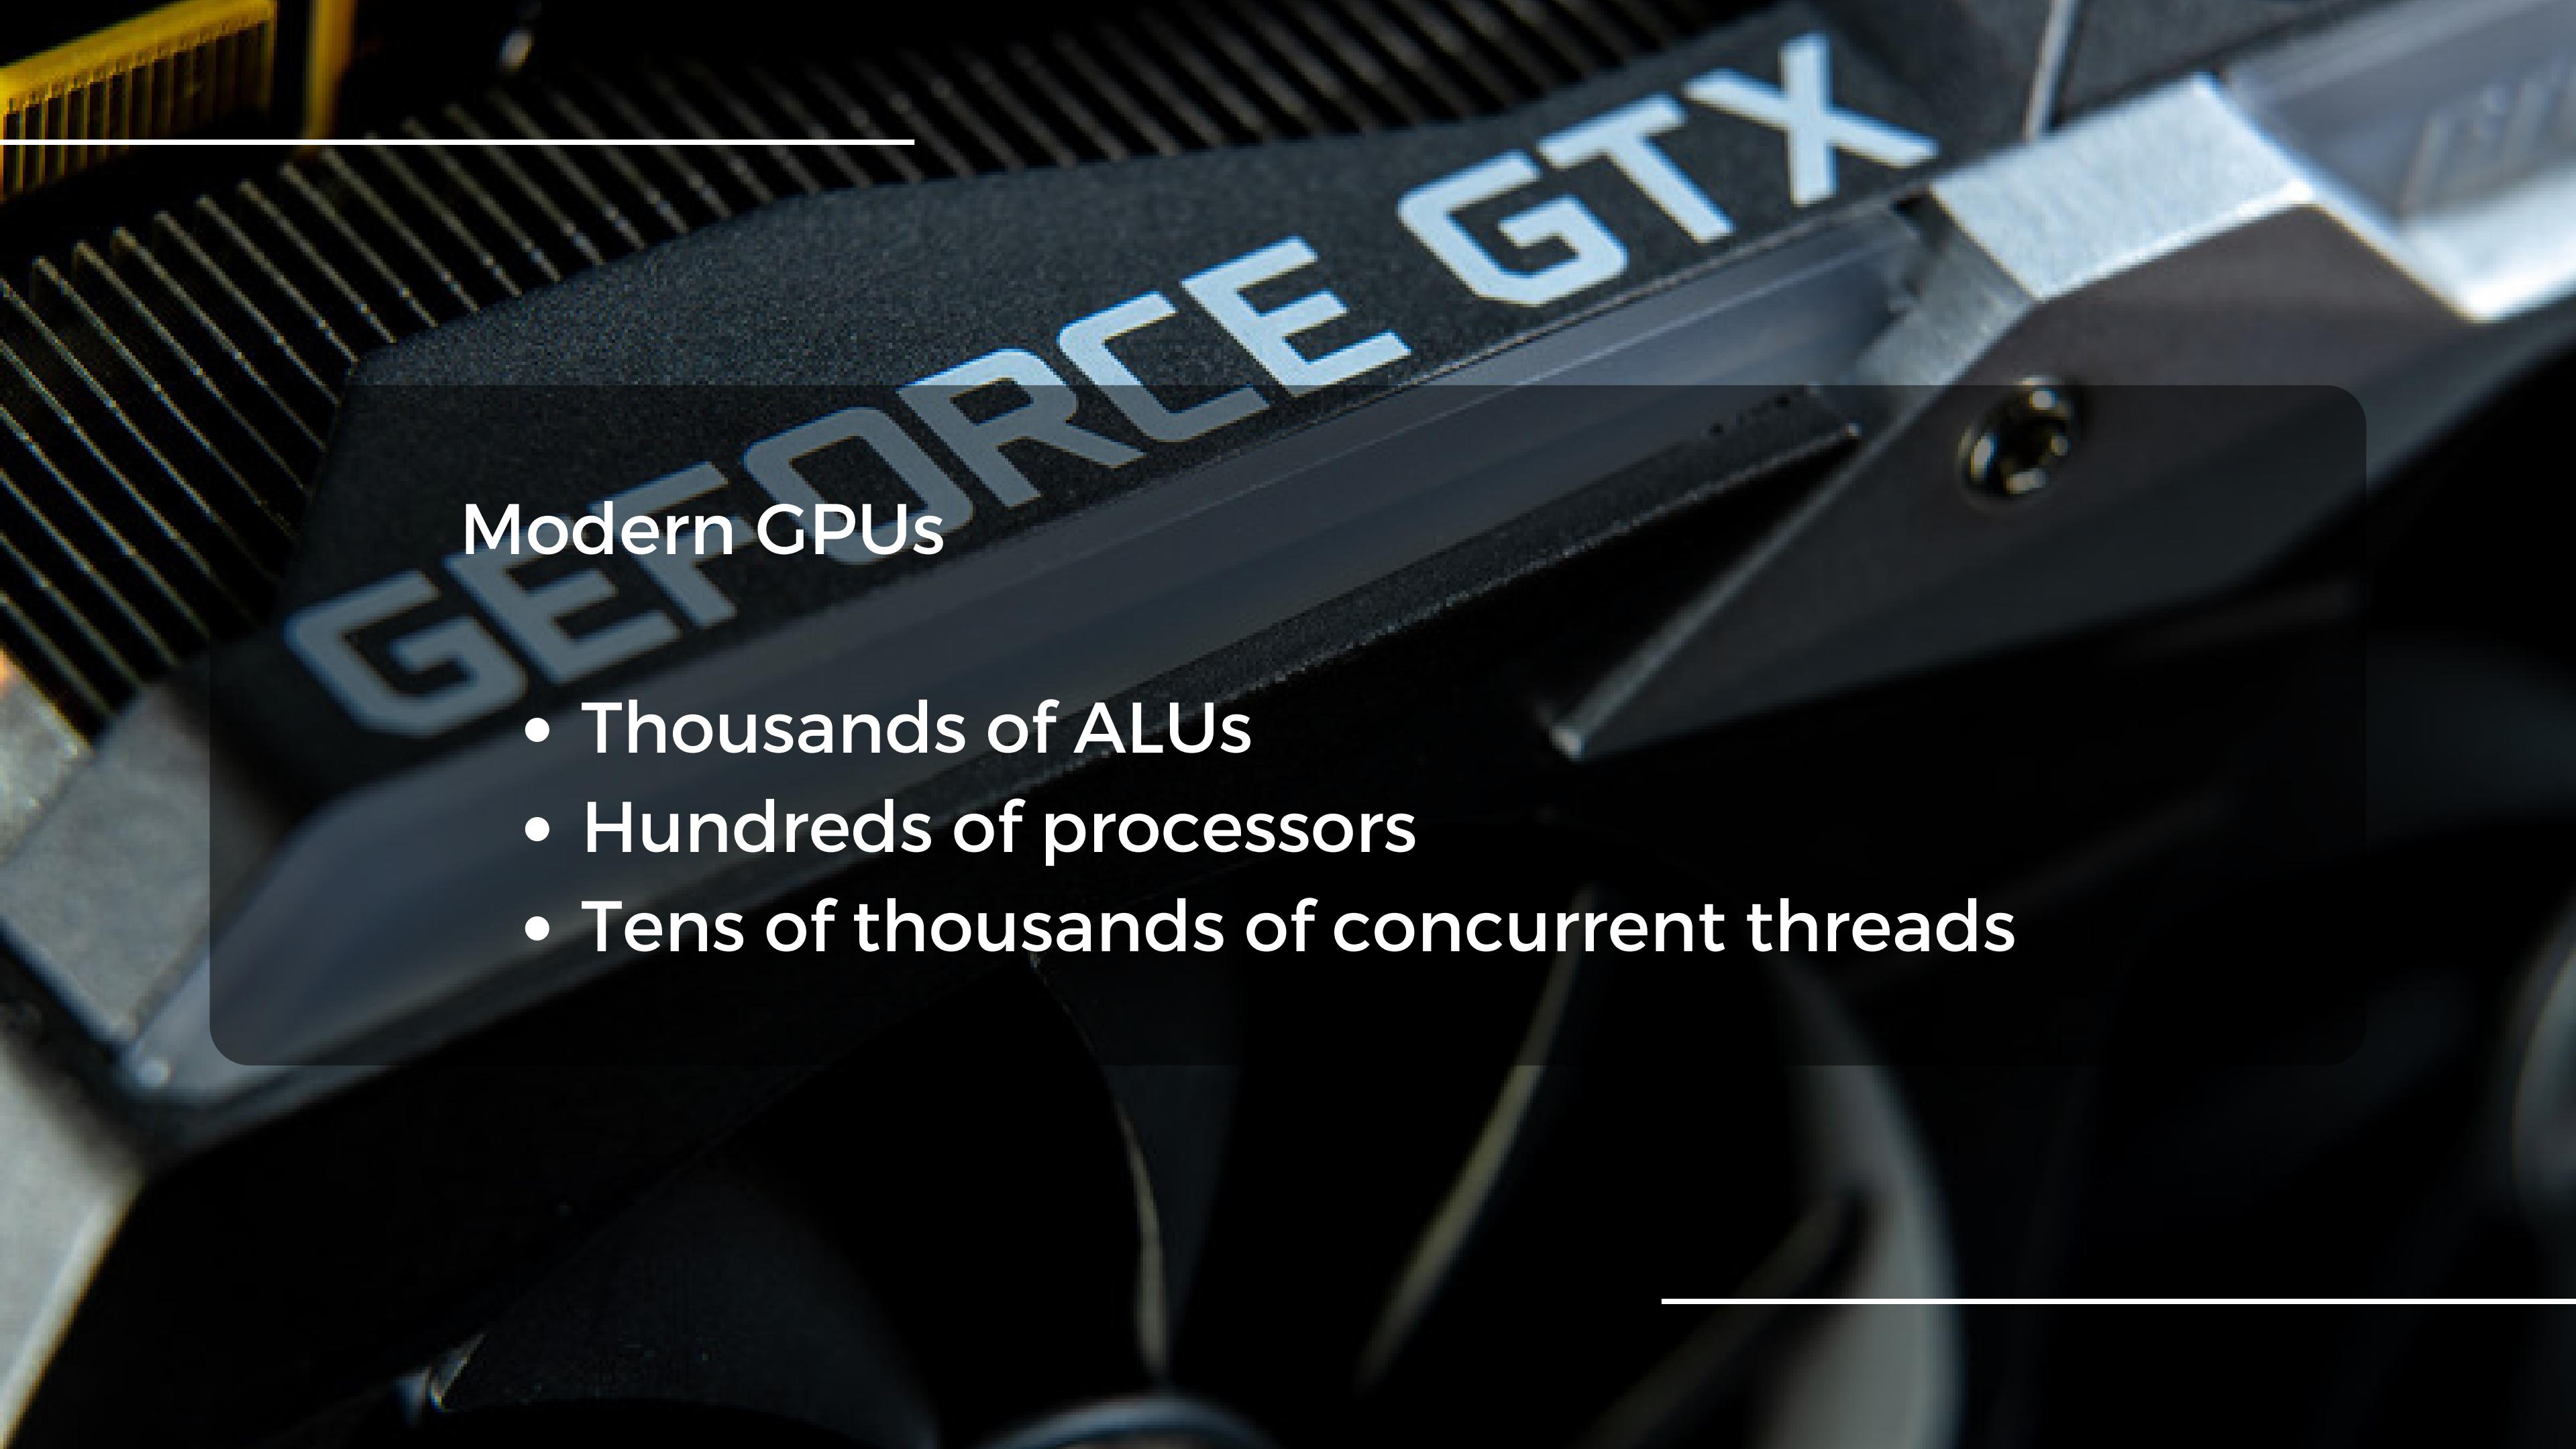
Seymour Cray



50 Years of Microprocessor Trend Data



Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2021 by K. Rupp



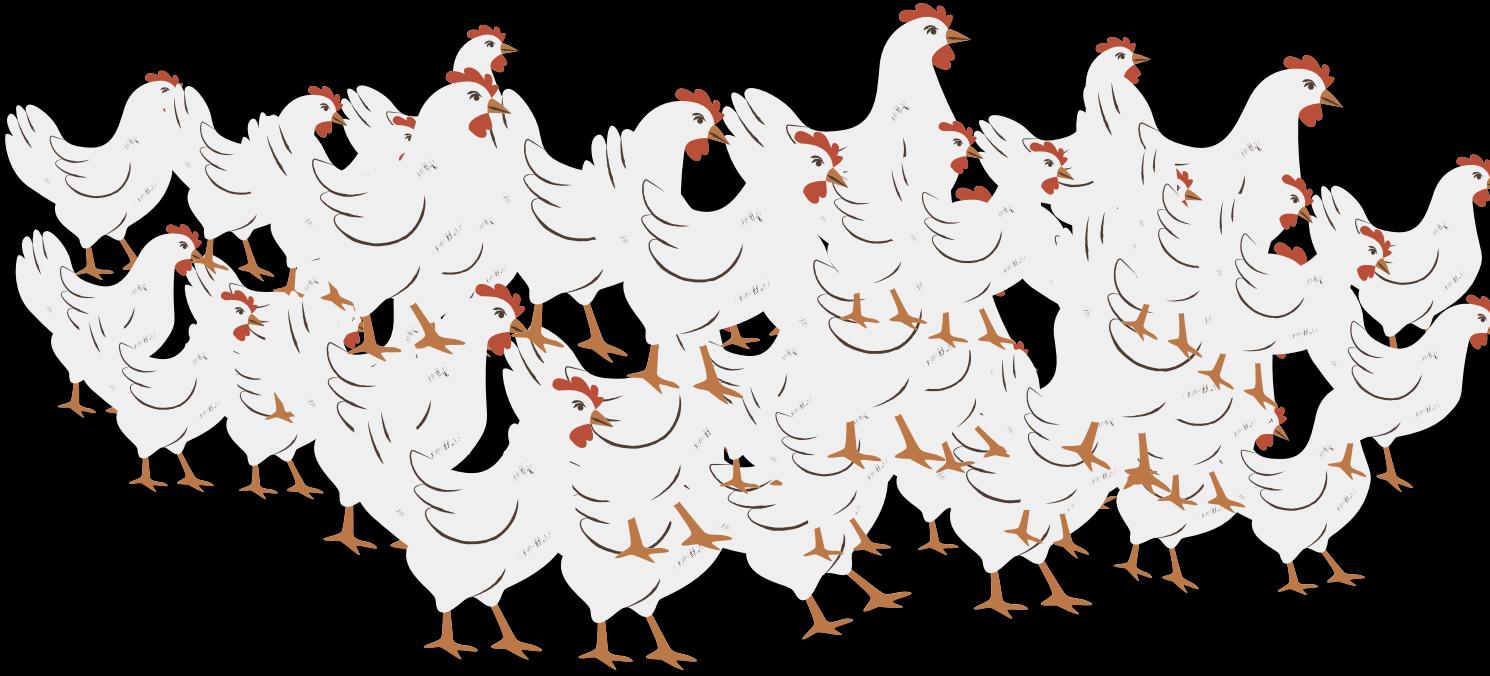
Modern GPUs

- Thousands of ALUs
- Hundreds of processors
- Tens of thousands of concurrent threads

We are doom

Why?

Managing chickens is hard!



CHAPTER 1

LAYING BRICKS



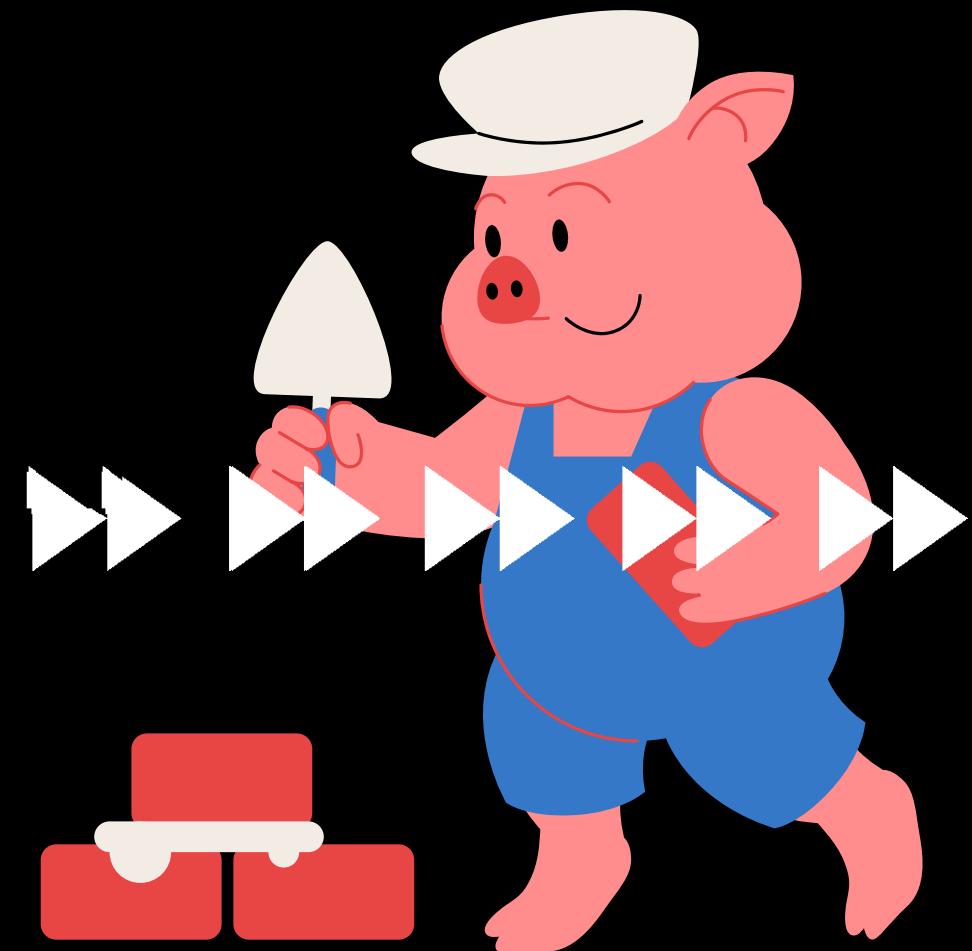
Just start laying bricks!

- Lay bricks 1 by 1
- By yourself



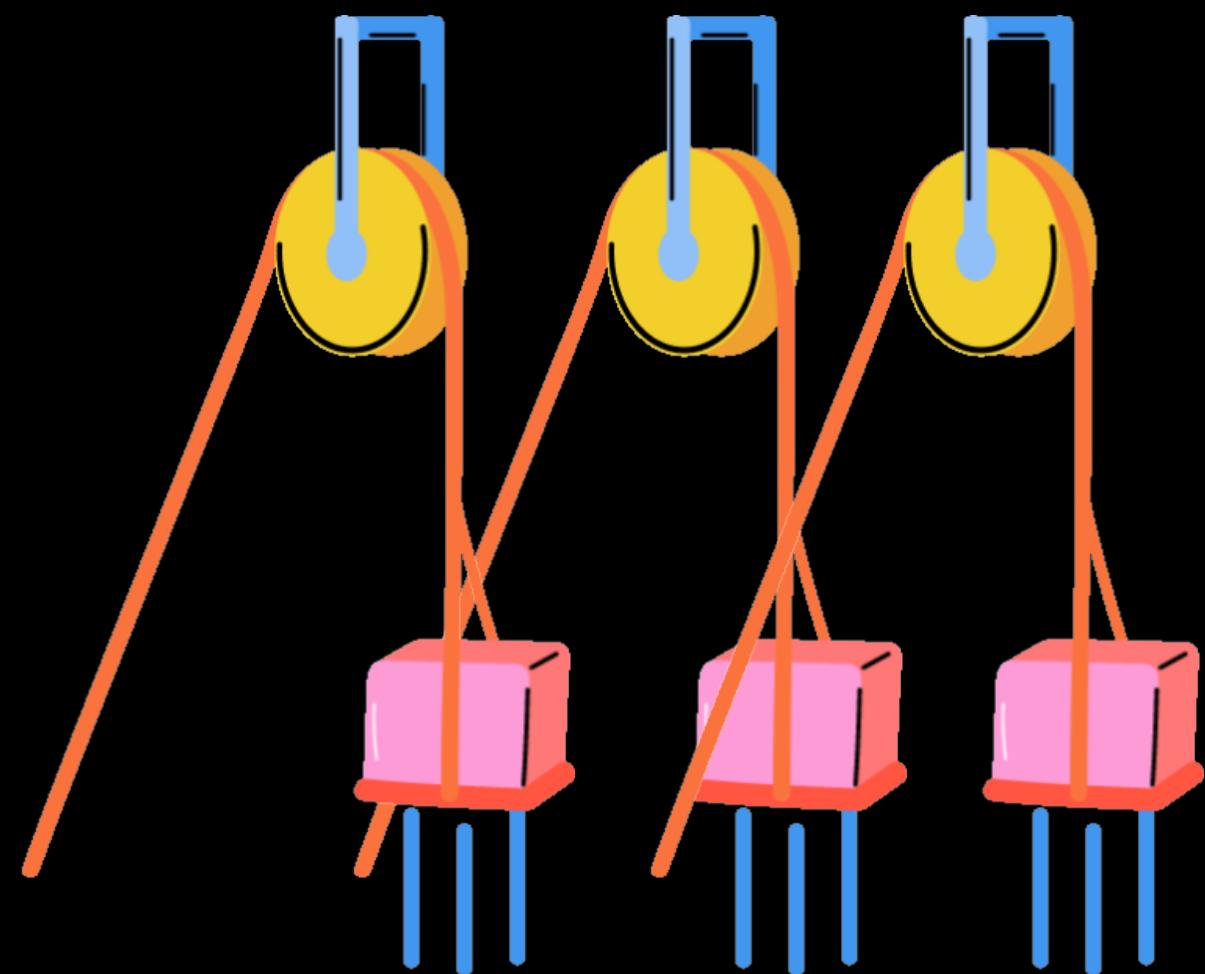
Lay faster!

- Instead of laying 5 brick per minute, lay 10 bricks per minute!
- There is a limit as to how fast a person lays brick.



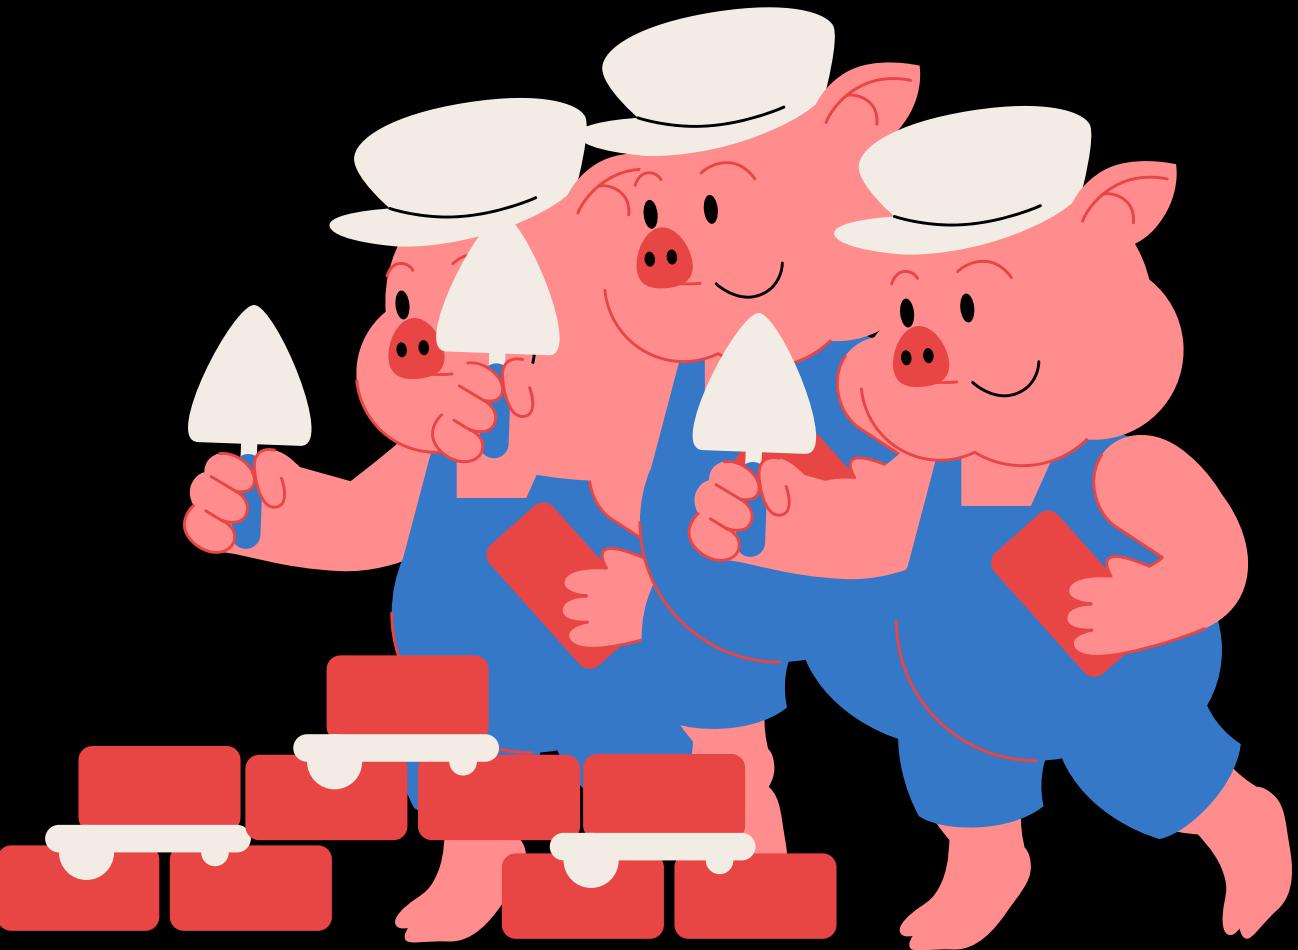
Use a machine 😱

- Buy a brick laying machine that can lay multiple bricks at once.



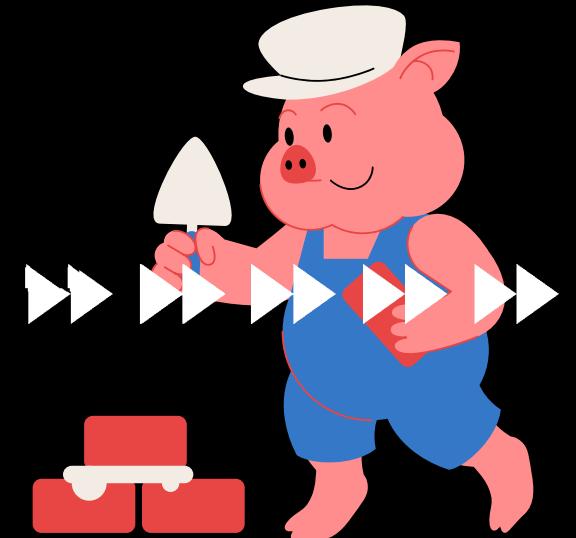
Hire more people

- Hire more people to lay bricks together!



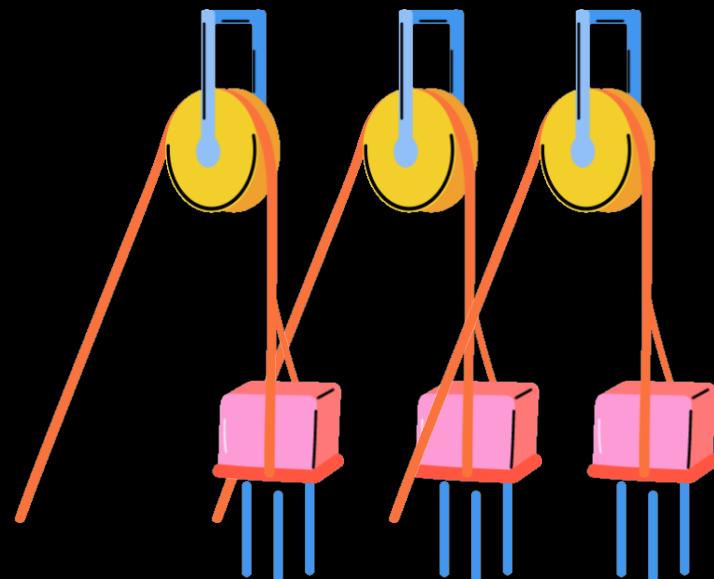
Analogy

Faster clock
cycles



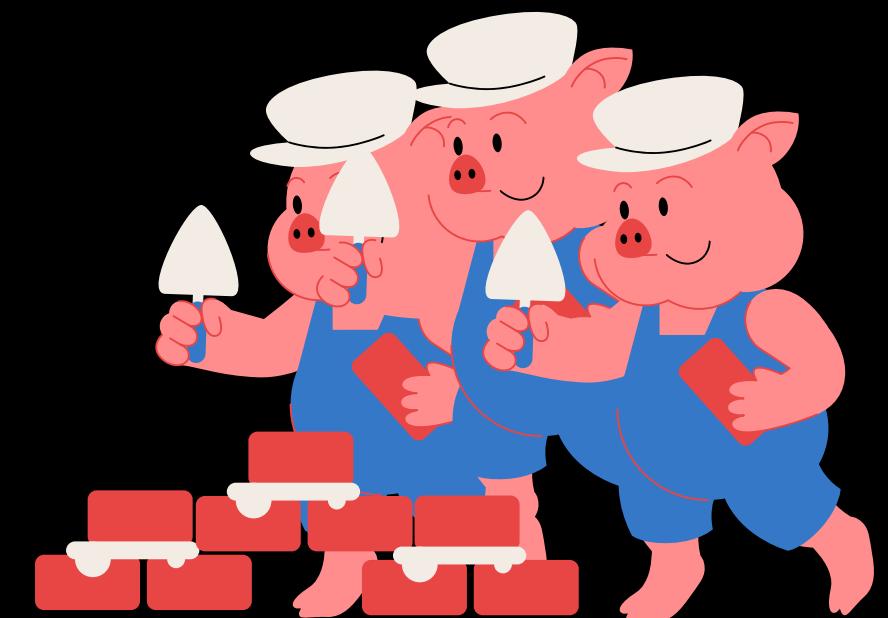
Lay faster

Instruction level
parallelism



Use a machine

More processors



Hire more people

CHAPTER 2

SOFTWARE PARALLELISM



Parallelism is a subset
of concurrency

**Parallelism is a subset
of **concurrency****

**Concurrency is dealing with a lot
of things at once. Parallelism is
doing a lot of things at once.**

- Rob Pike

What is software parallelism?

- Divide the work into units of tasks.
- Run independently.
- No dependencies from other tasks.
- Can be executed simultaneously.

How it ends up being executed on the hardware is completely irrelevant to us.

CHAPTER 3

DATA-PARALLEL PRIMITIVES



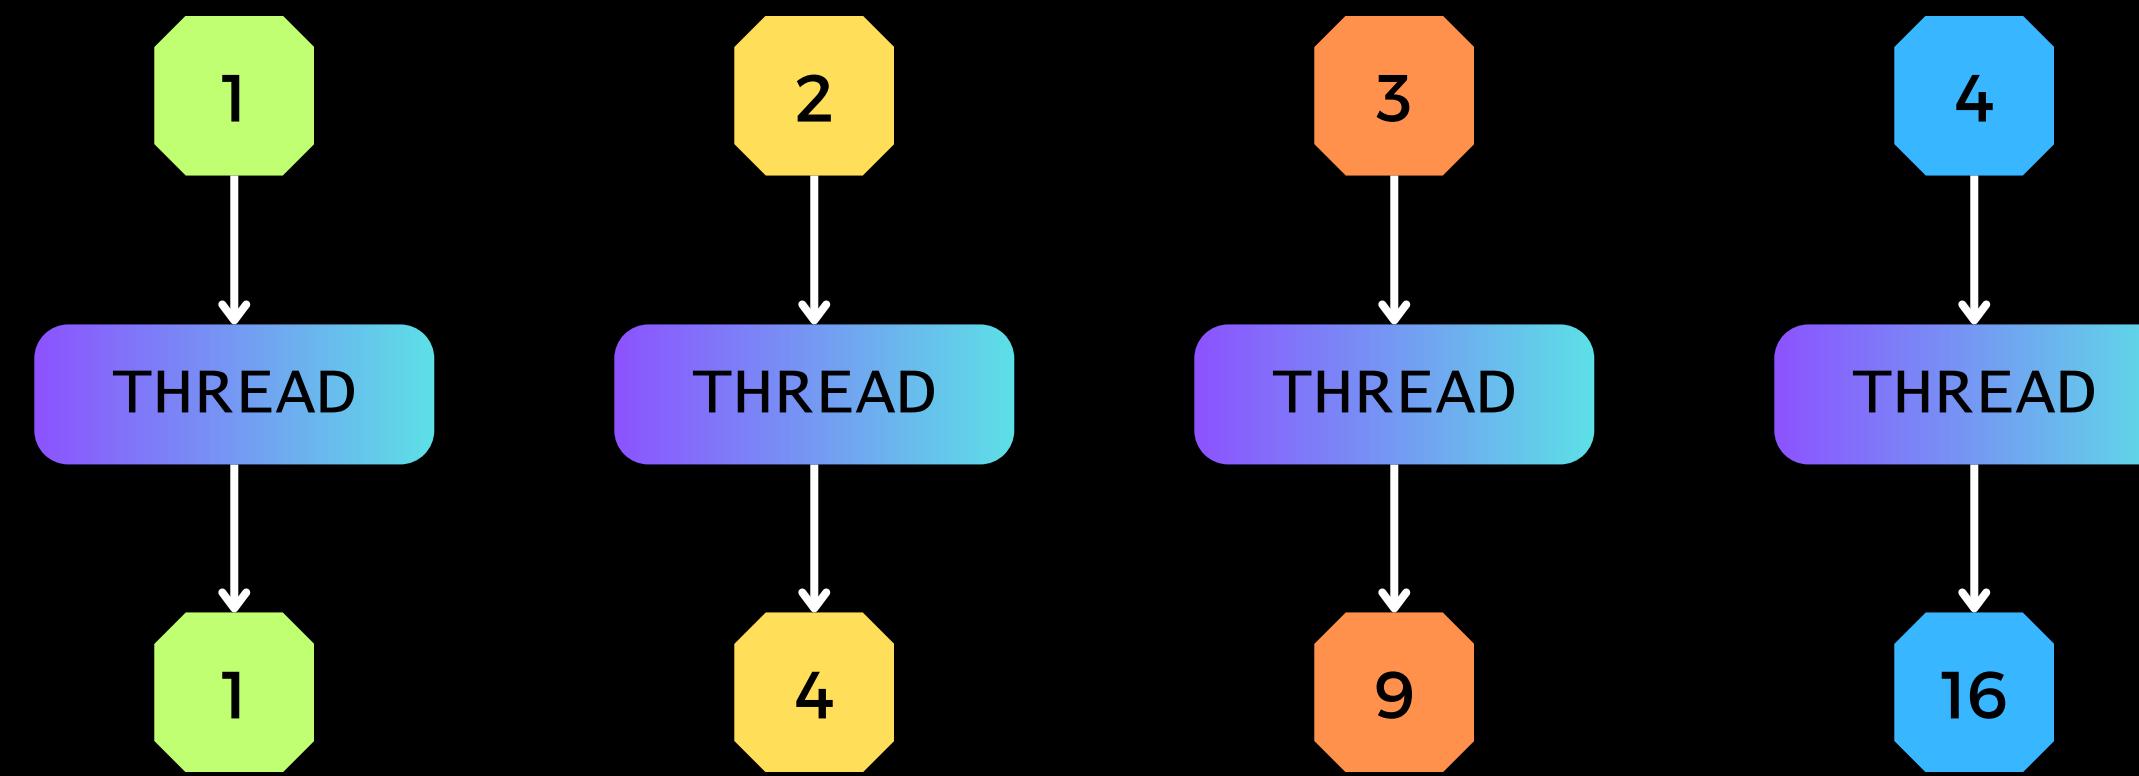
What are data-parallel primitives? (DPP)

- Building blocks.
 - Commonly used patterns.
 - Combined to create complex parallel algorithms.
 - Reuse existing high performance implementations.
-

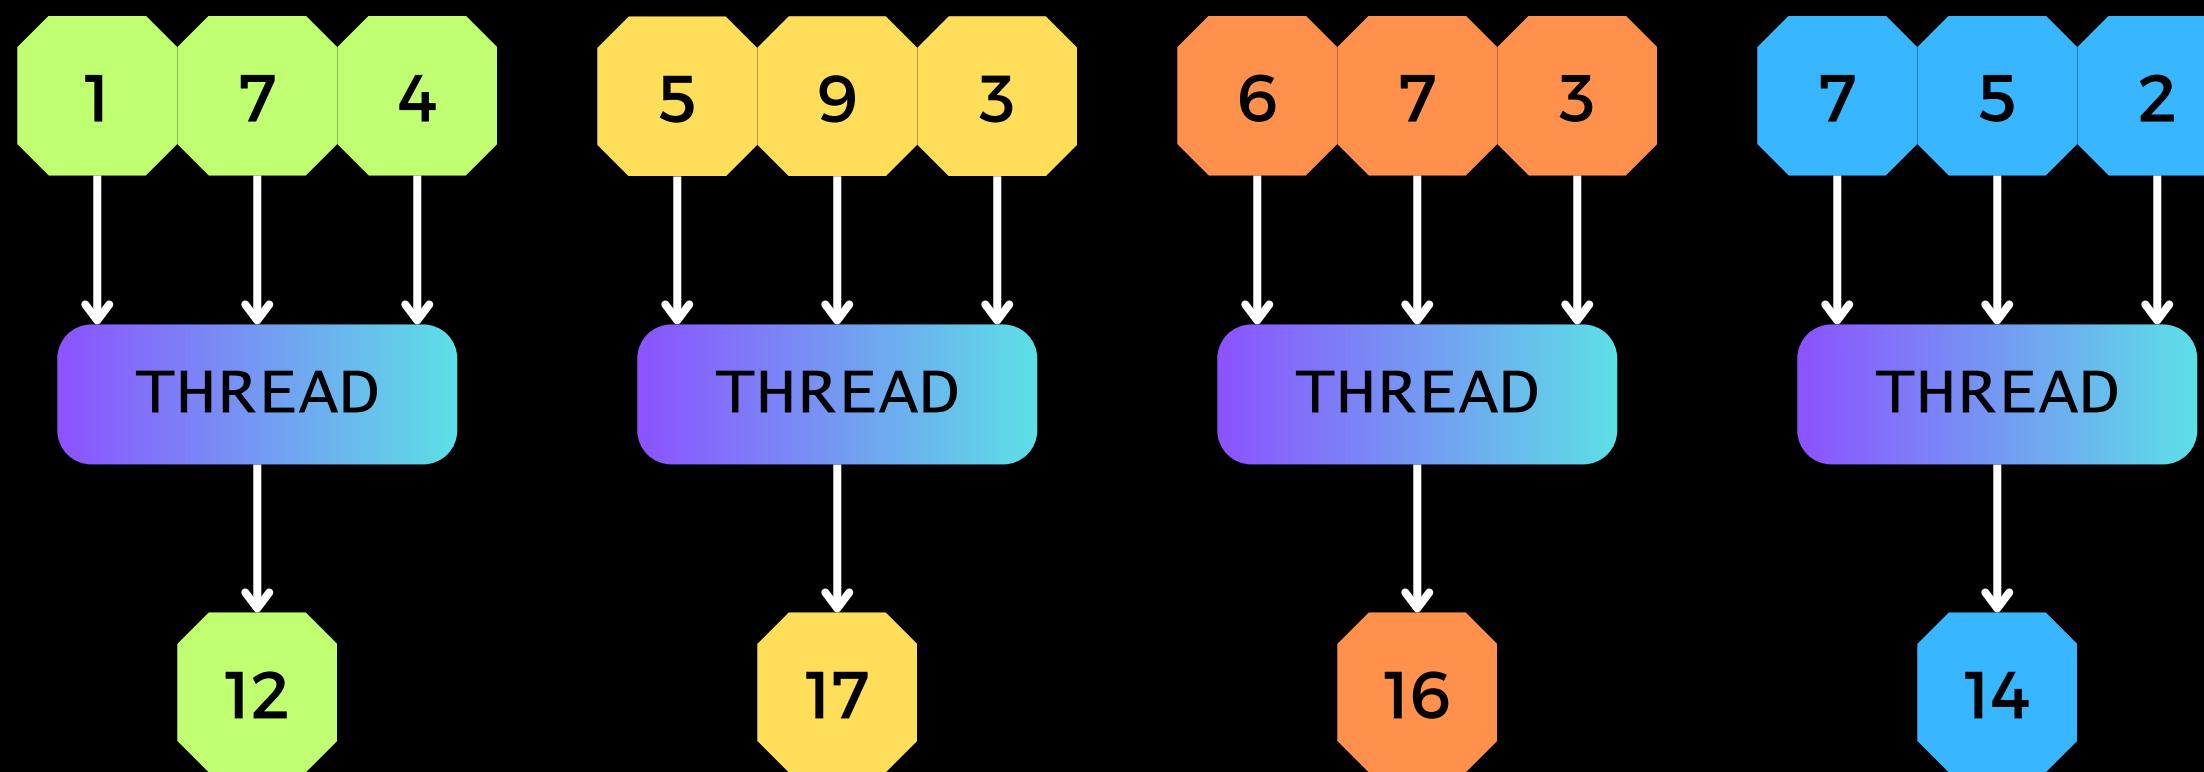
Common Types of DPP

- Map
 - Stencil
 - Reduce
 - Scan
 - Gather/scatter
 - Compaction
 - Partition
-

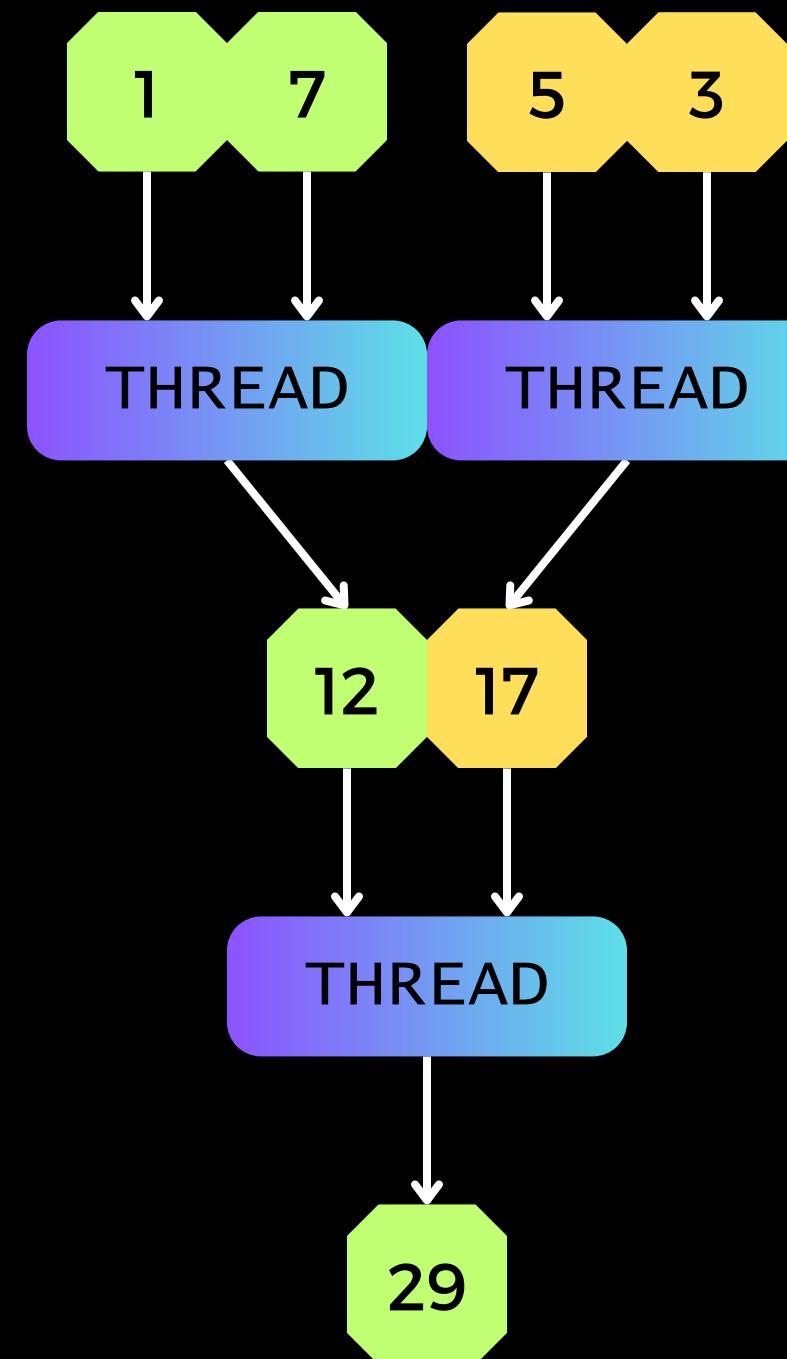
Map



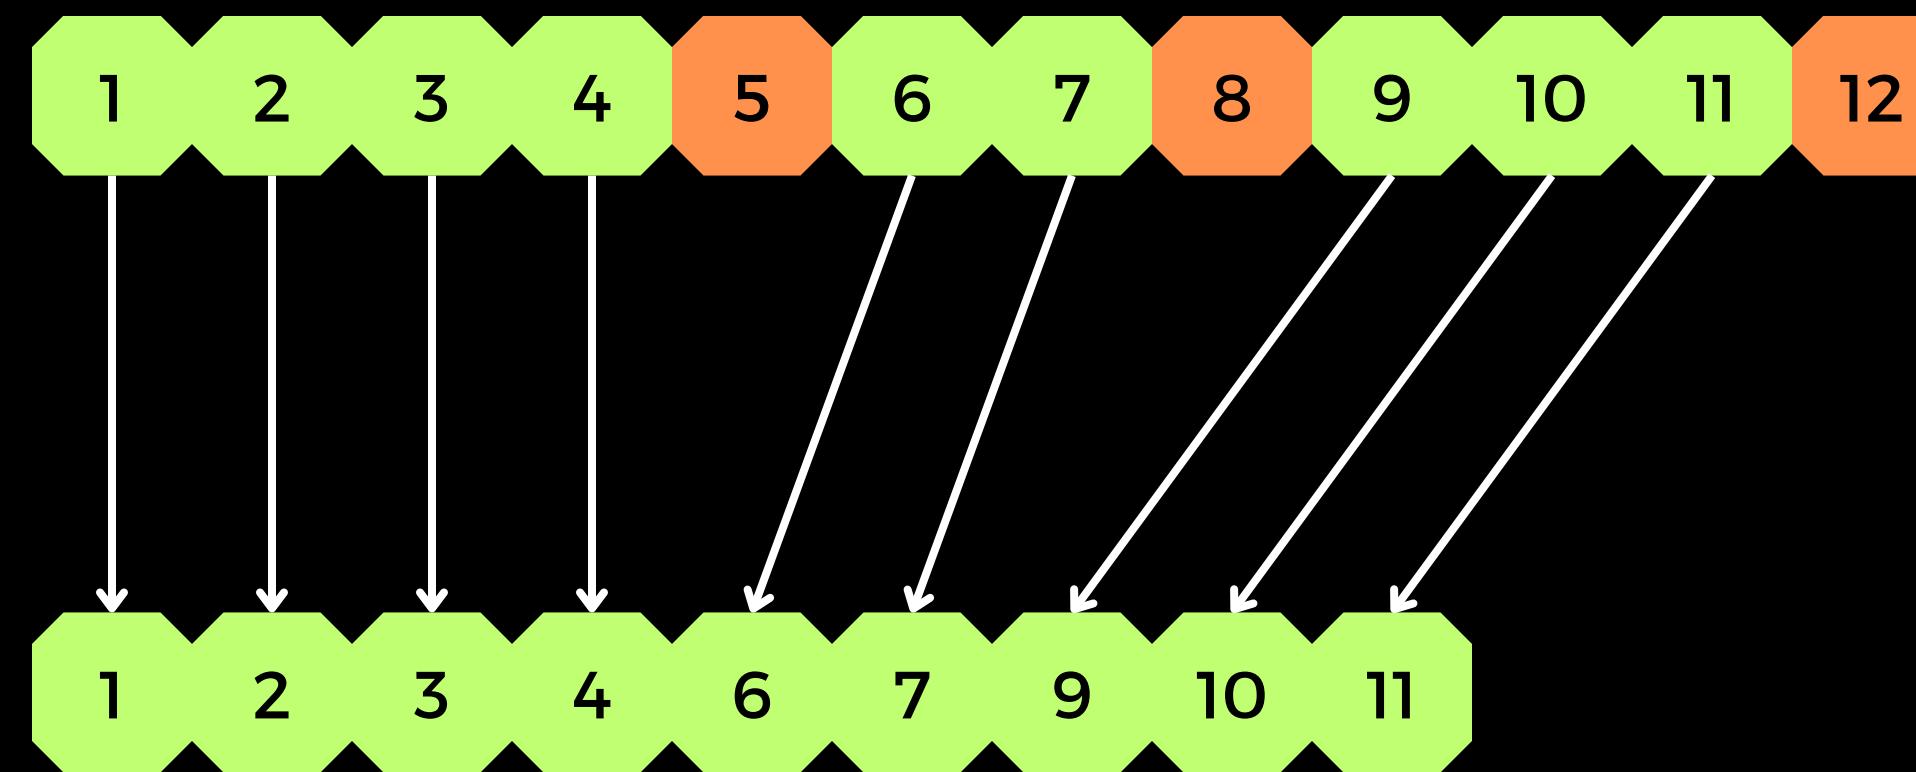
Stencil



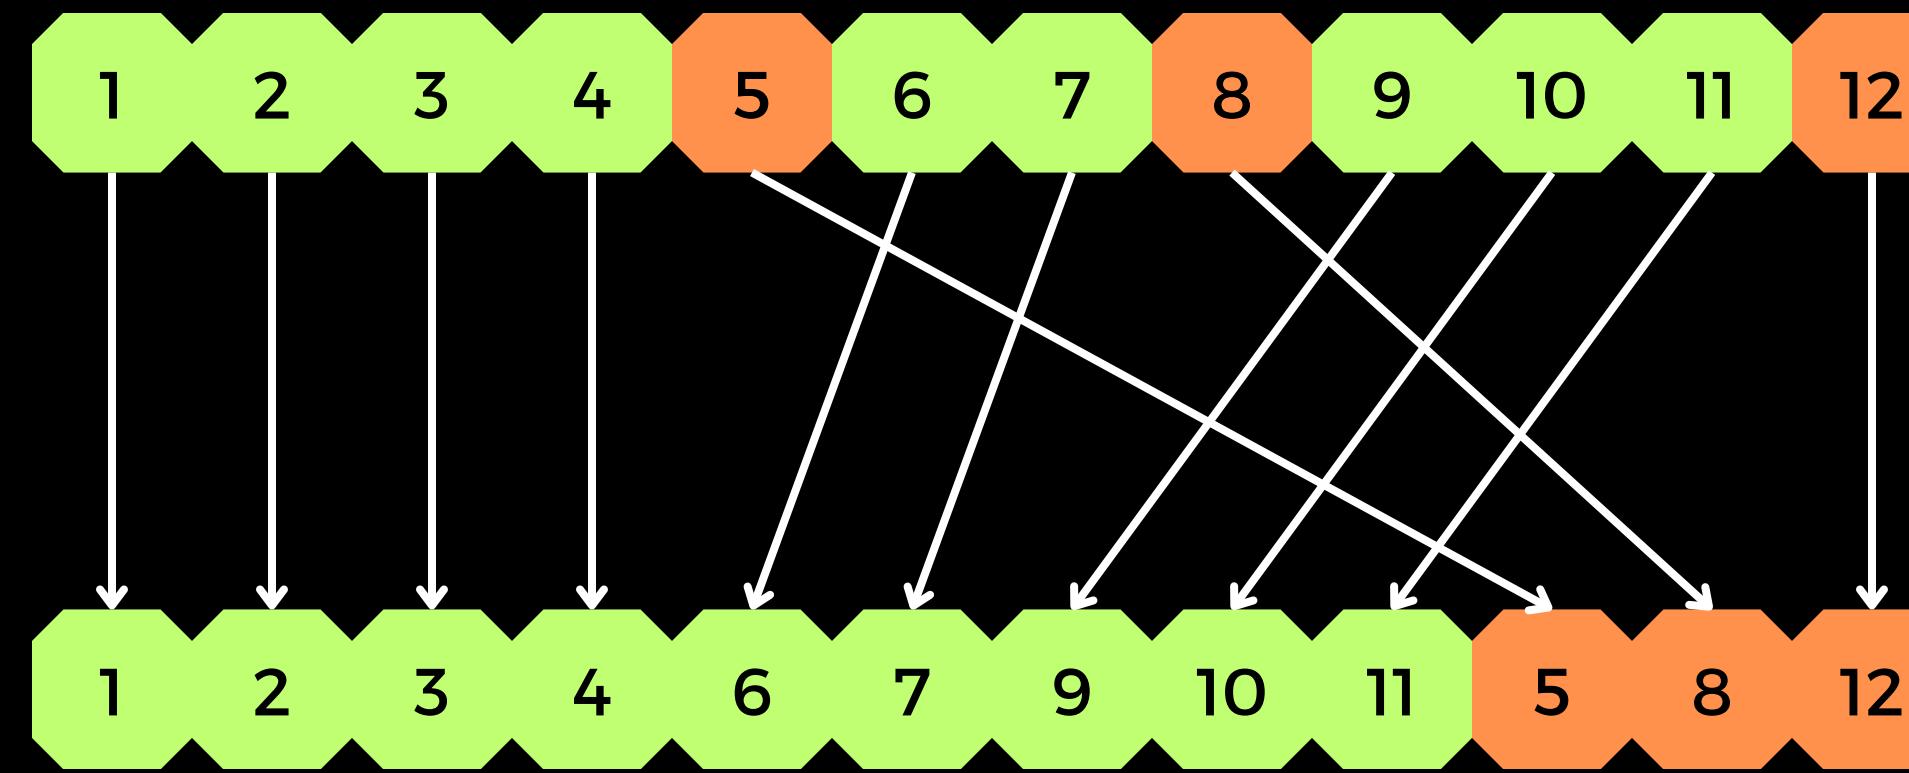
Reduce



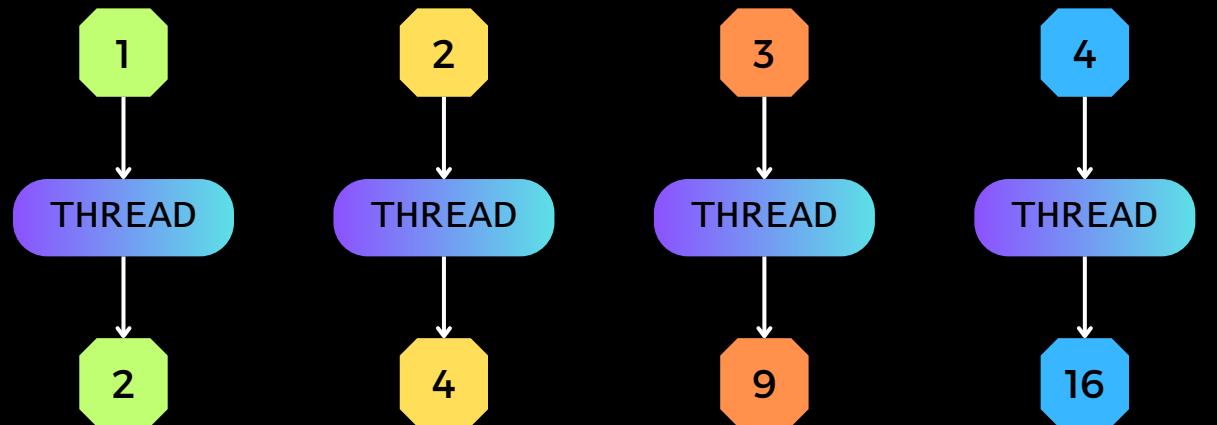
Compaction



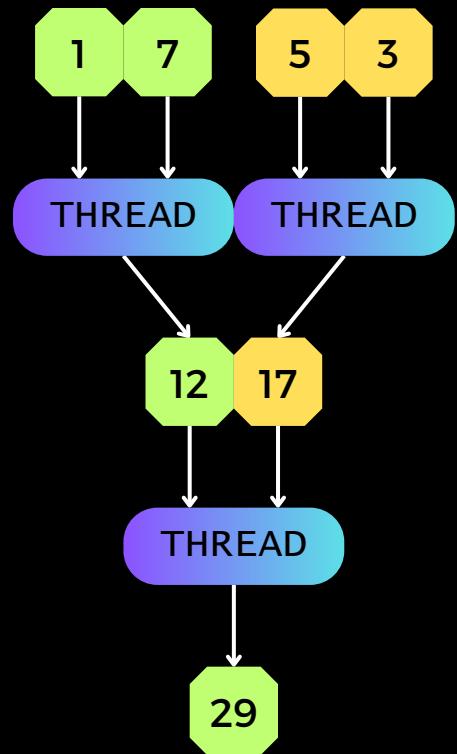
Partition



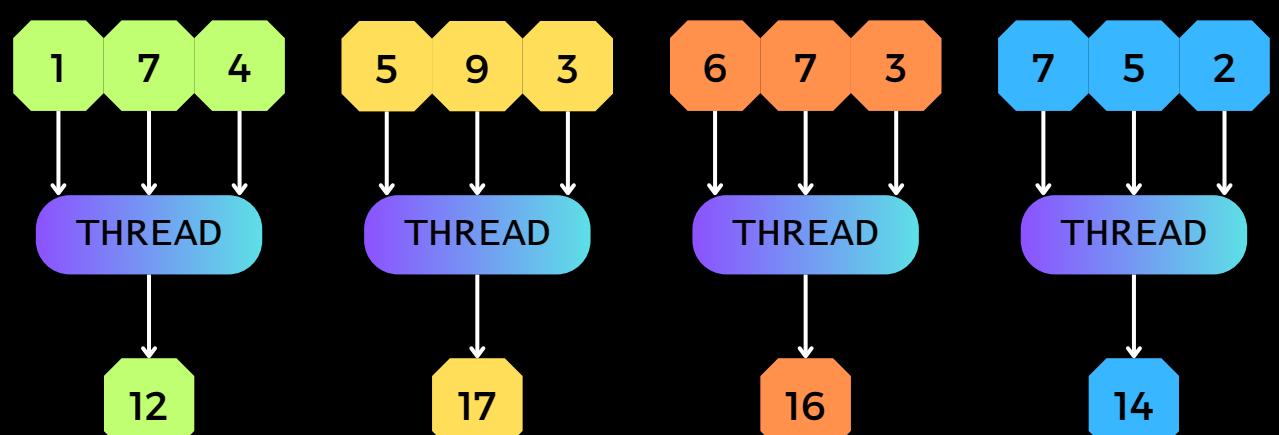
Map



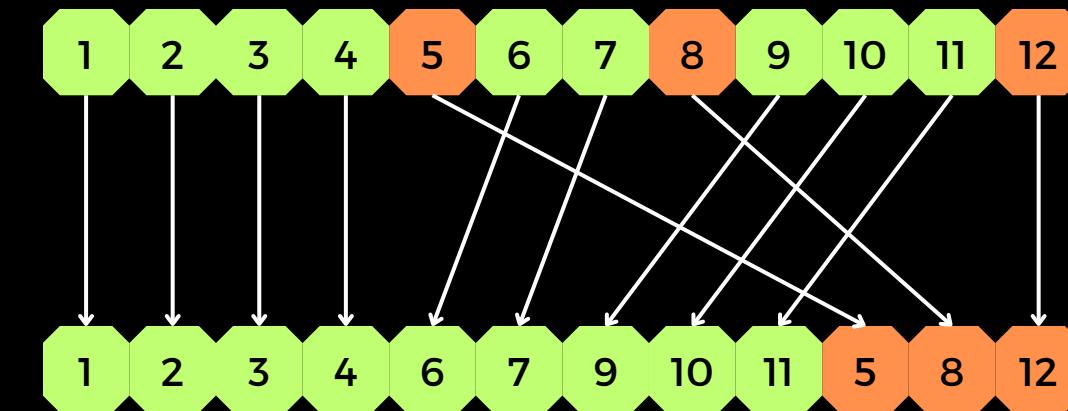
Reduce



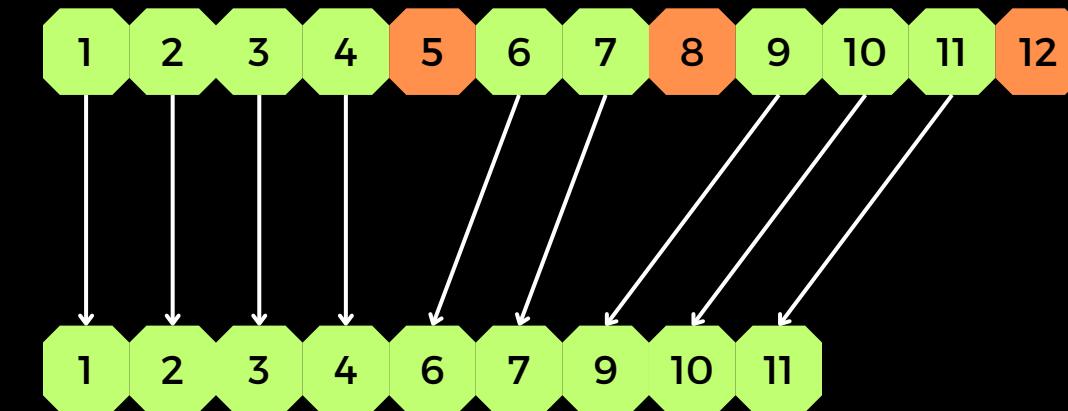
Stencil



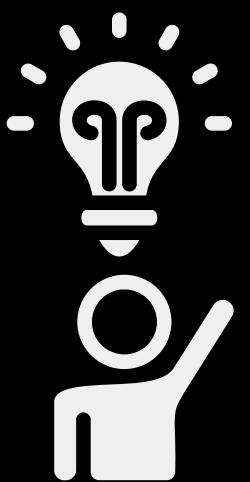
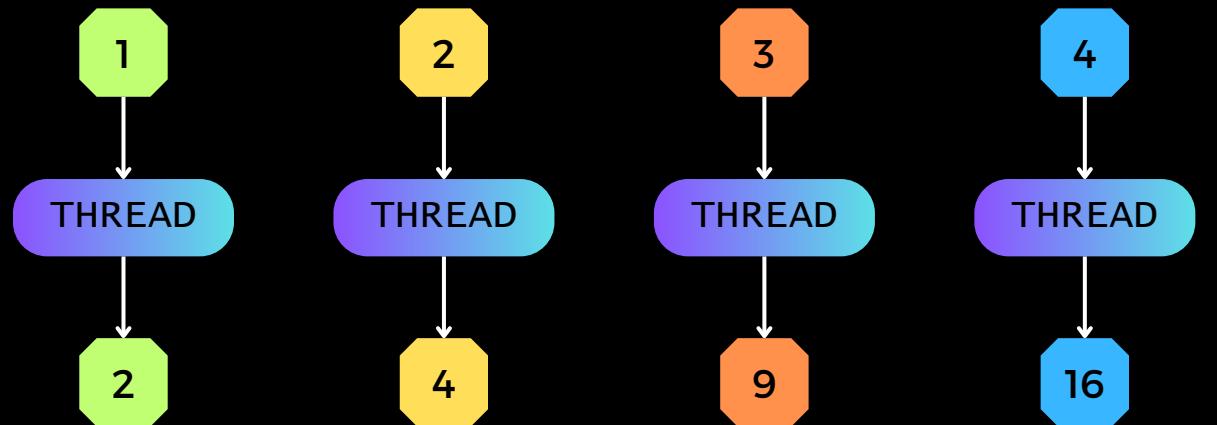
Partition



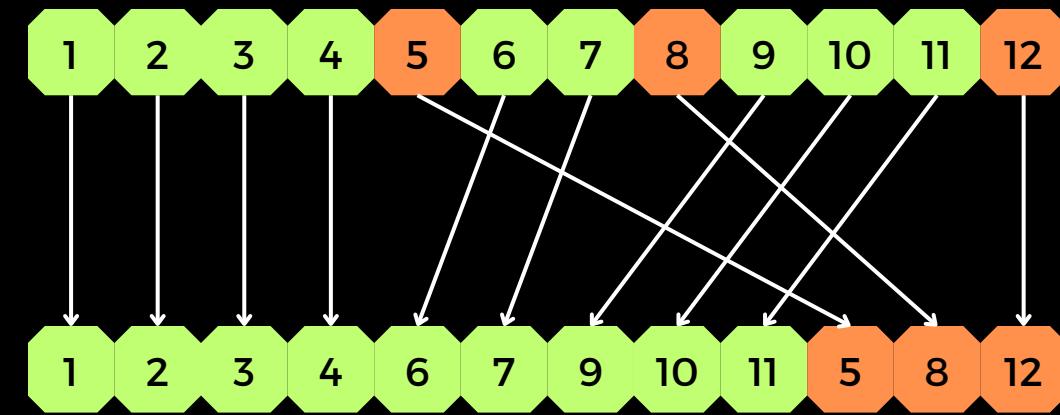
Compaction



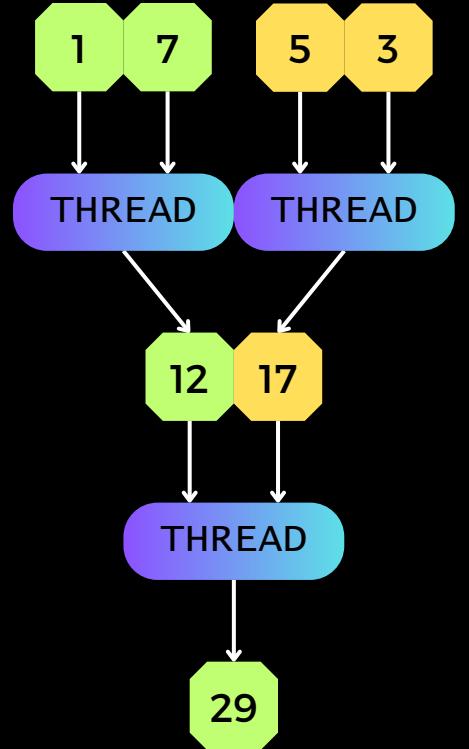
Map



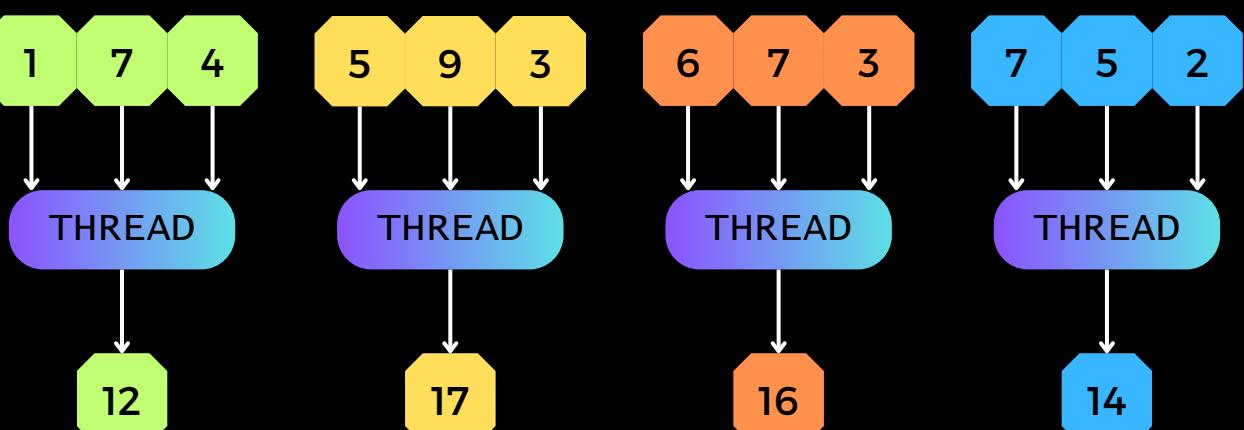
Partition



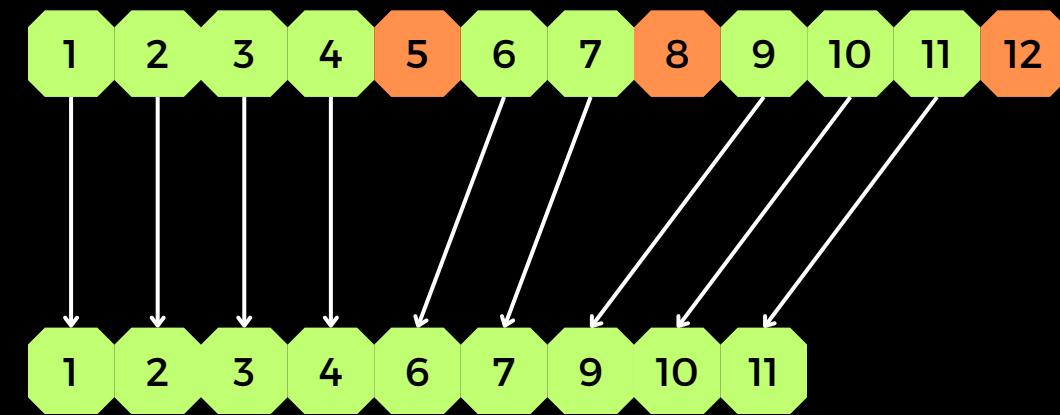
Reduce



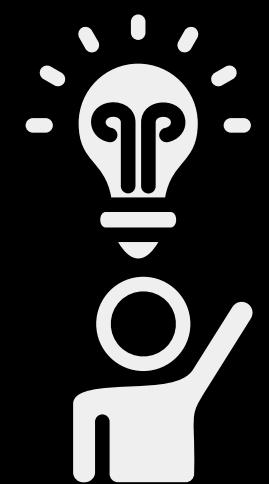
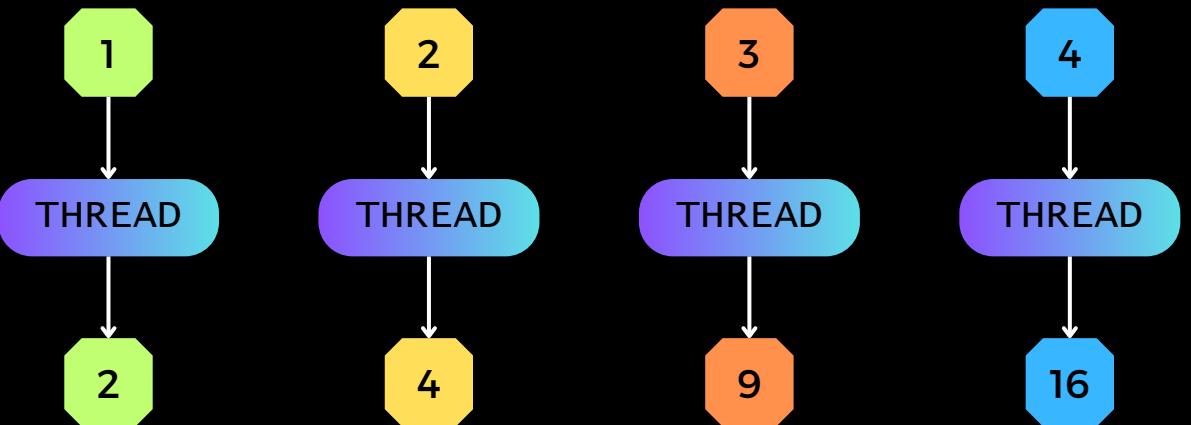
Stencil



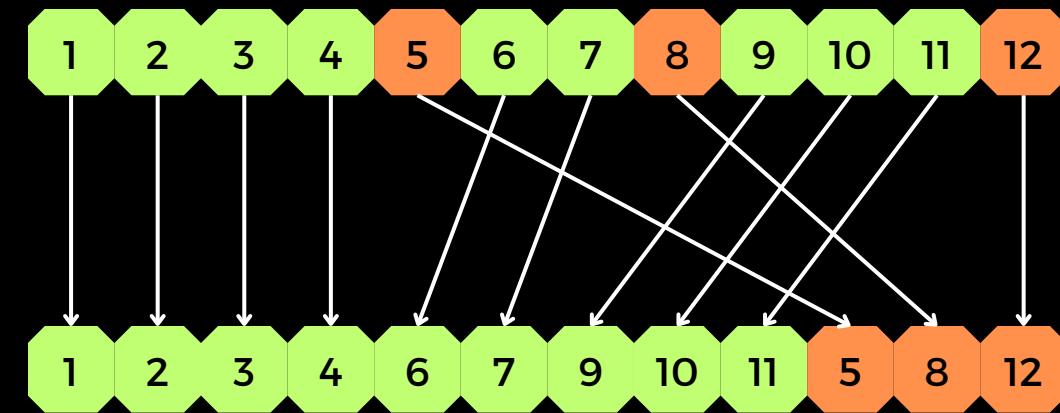
Compaction



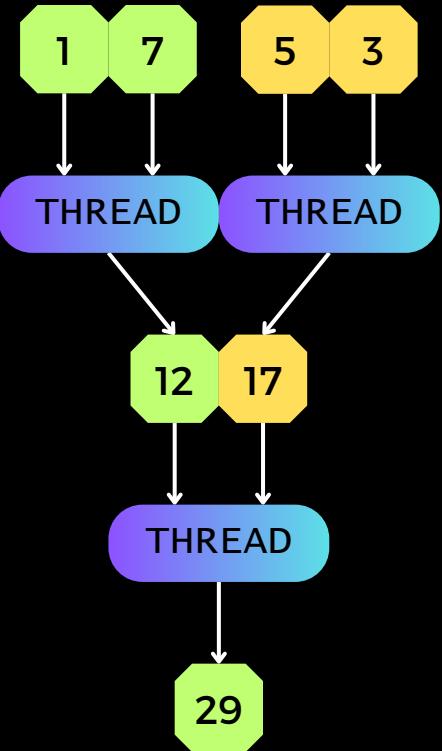
Map



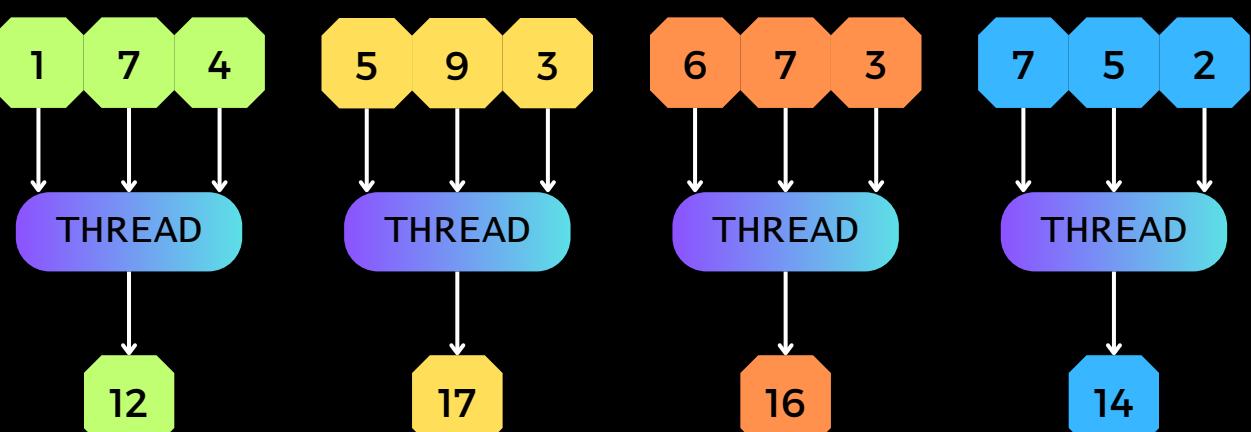
Partition



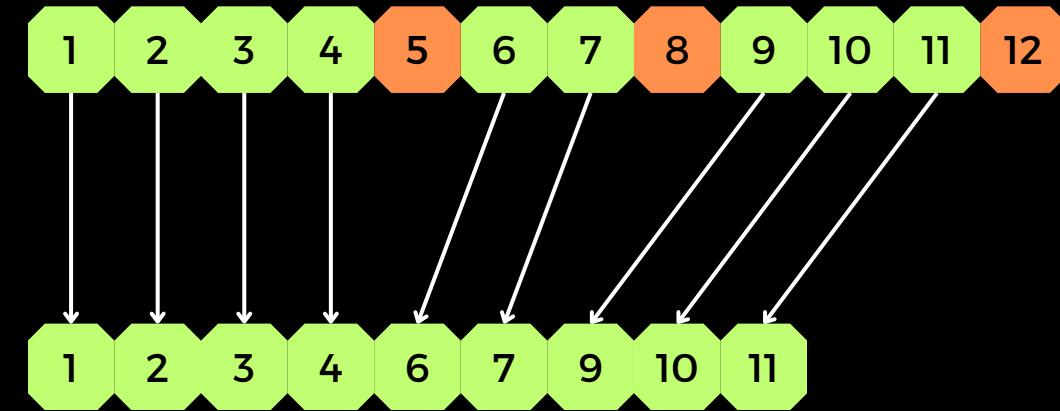
Reduce



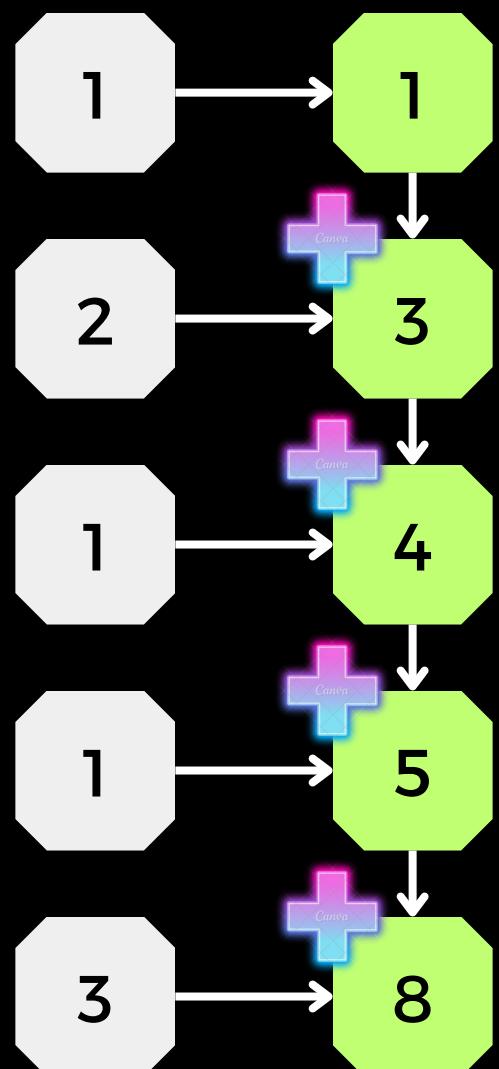
Stencil



Compaction

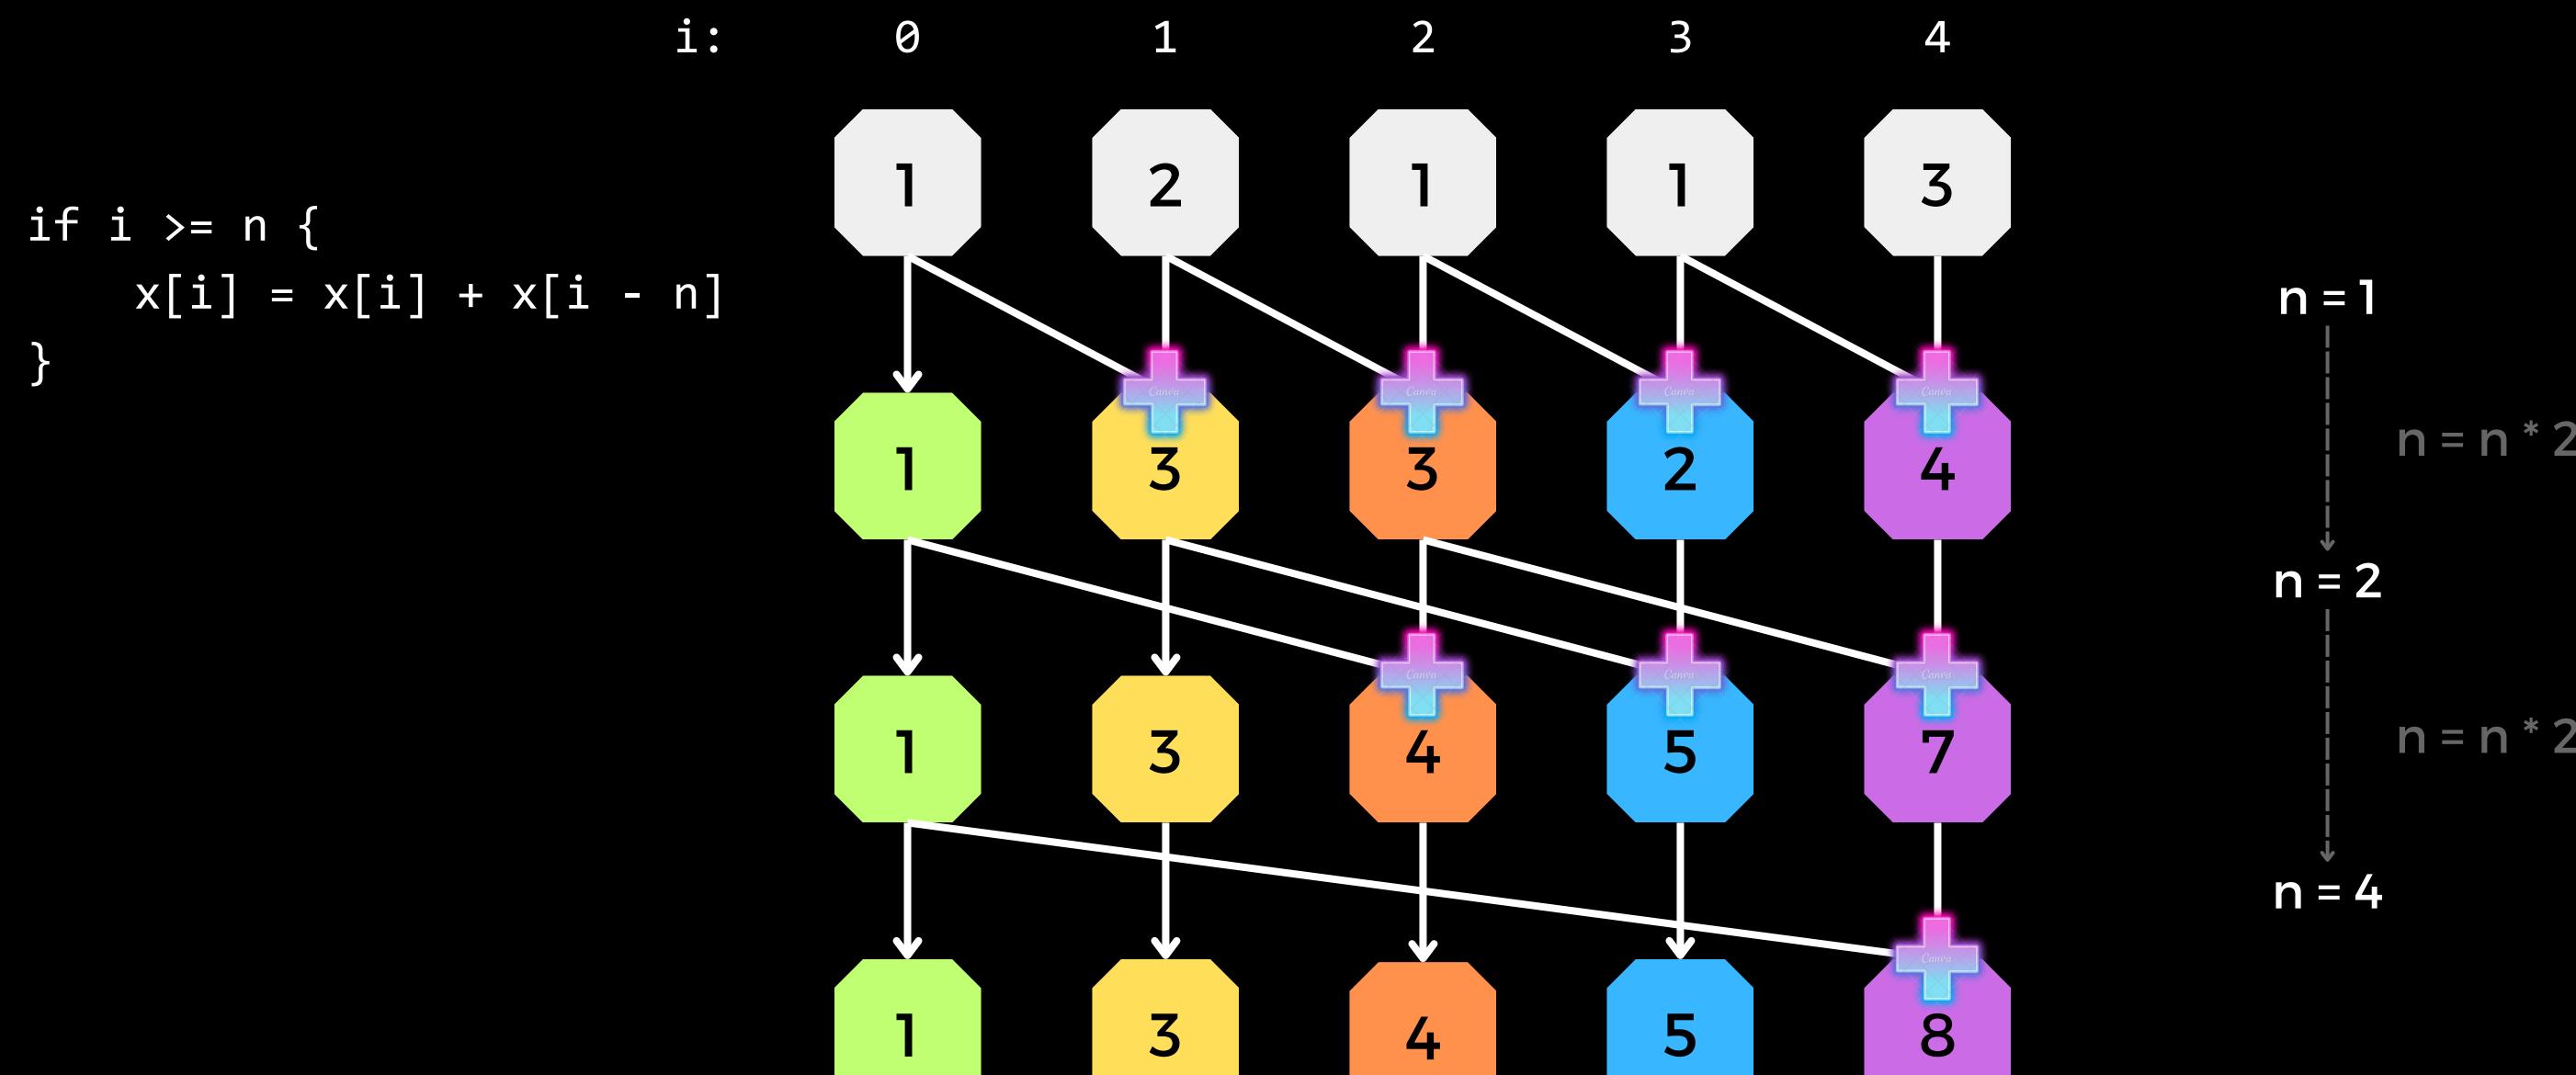


Sum Scan (in serial)



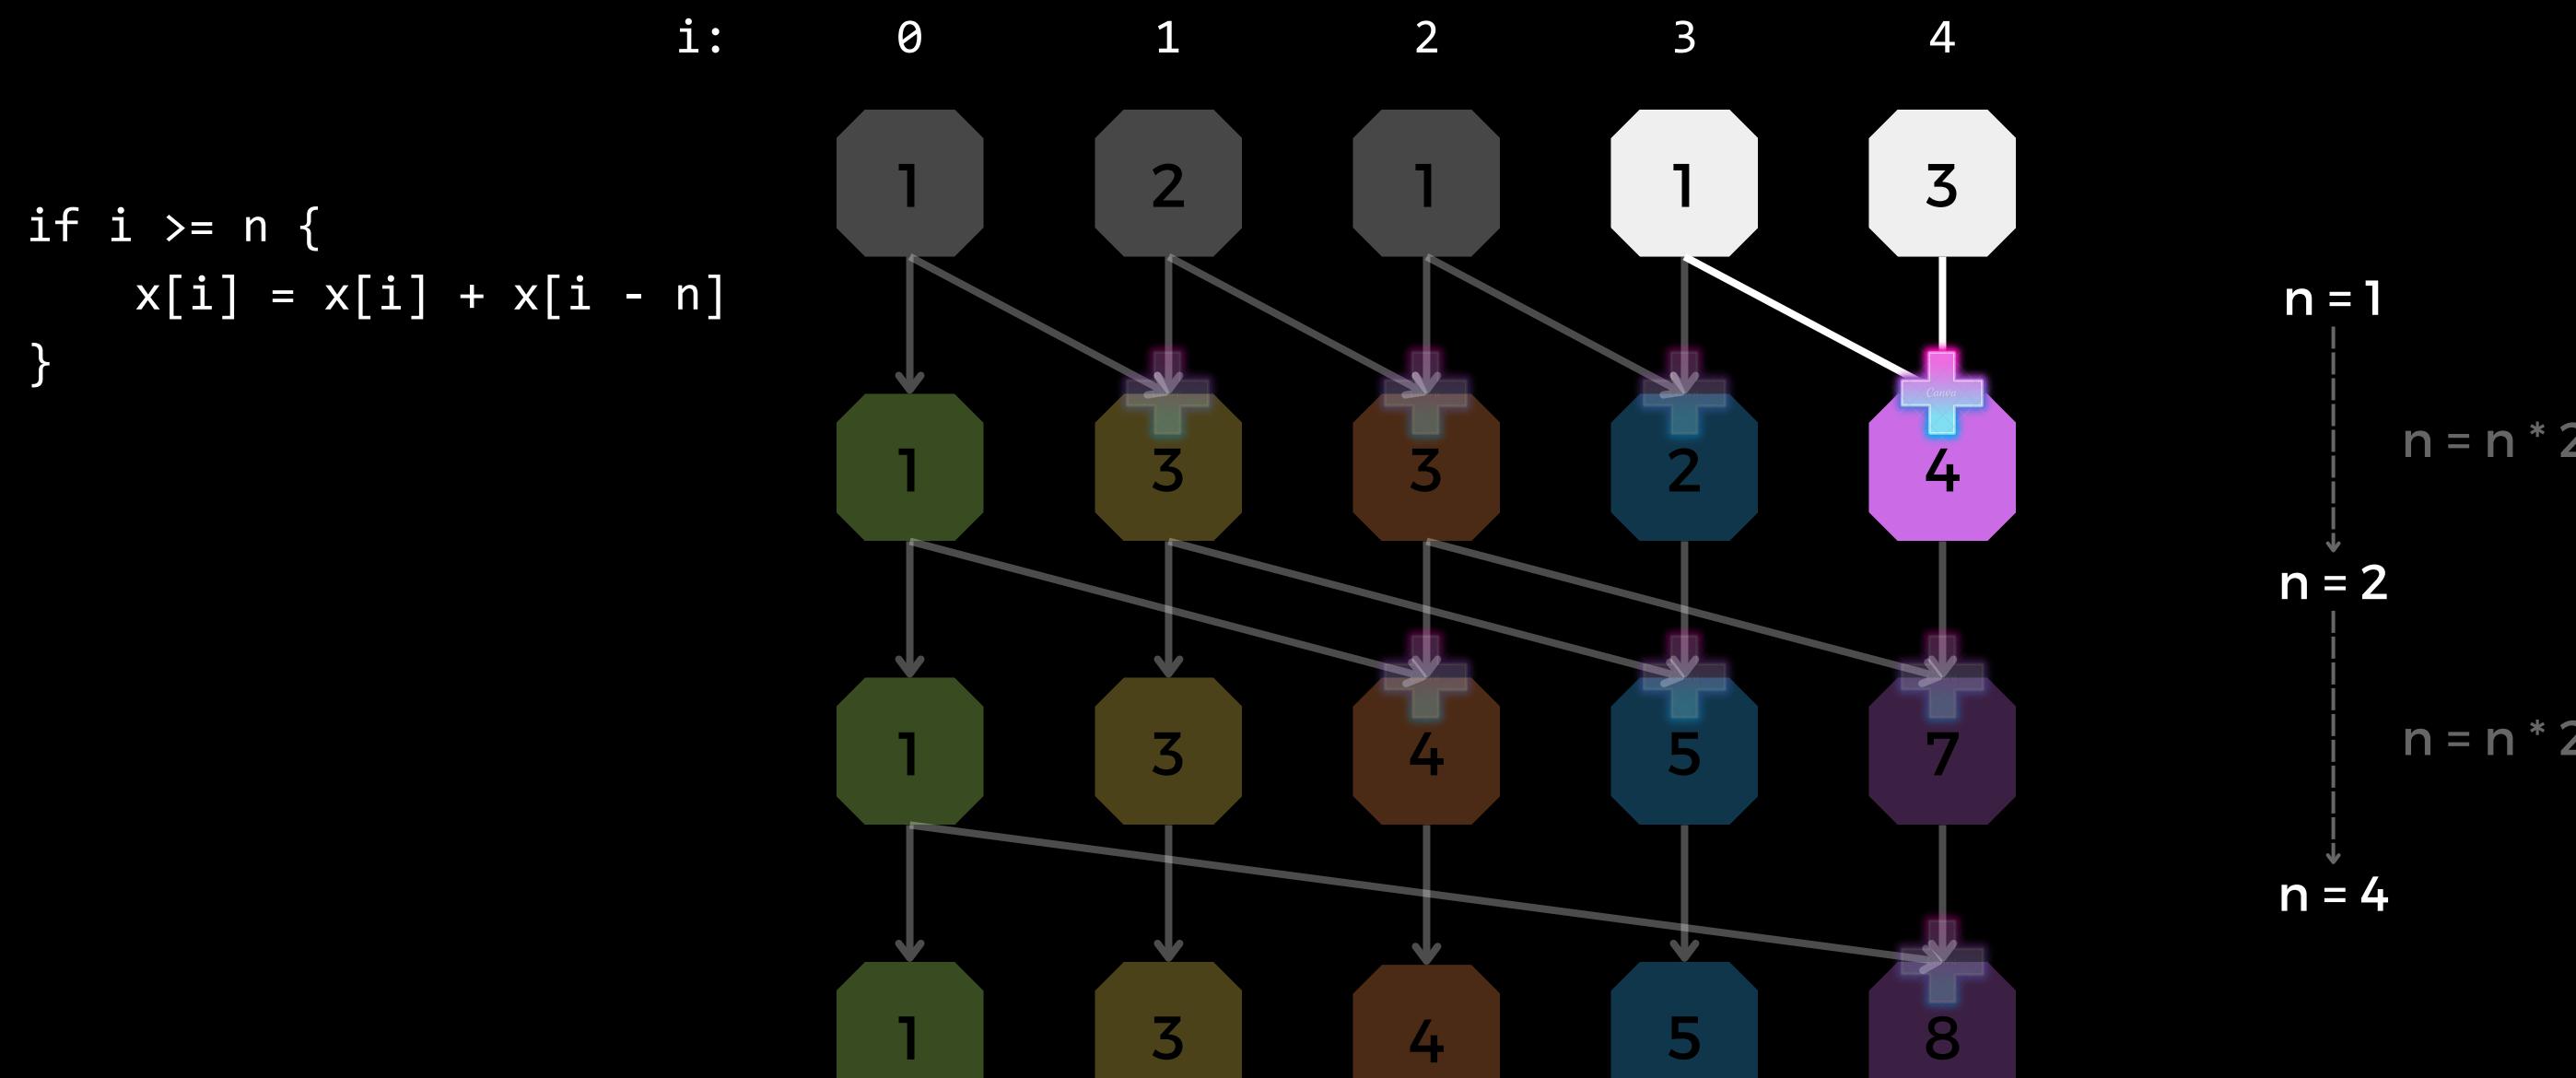
$O(n)$

Sum Scan



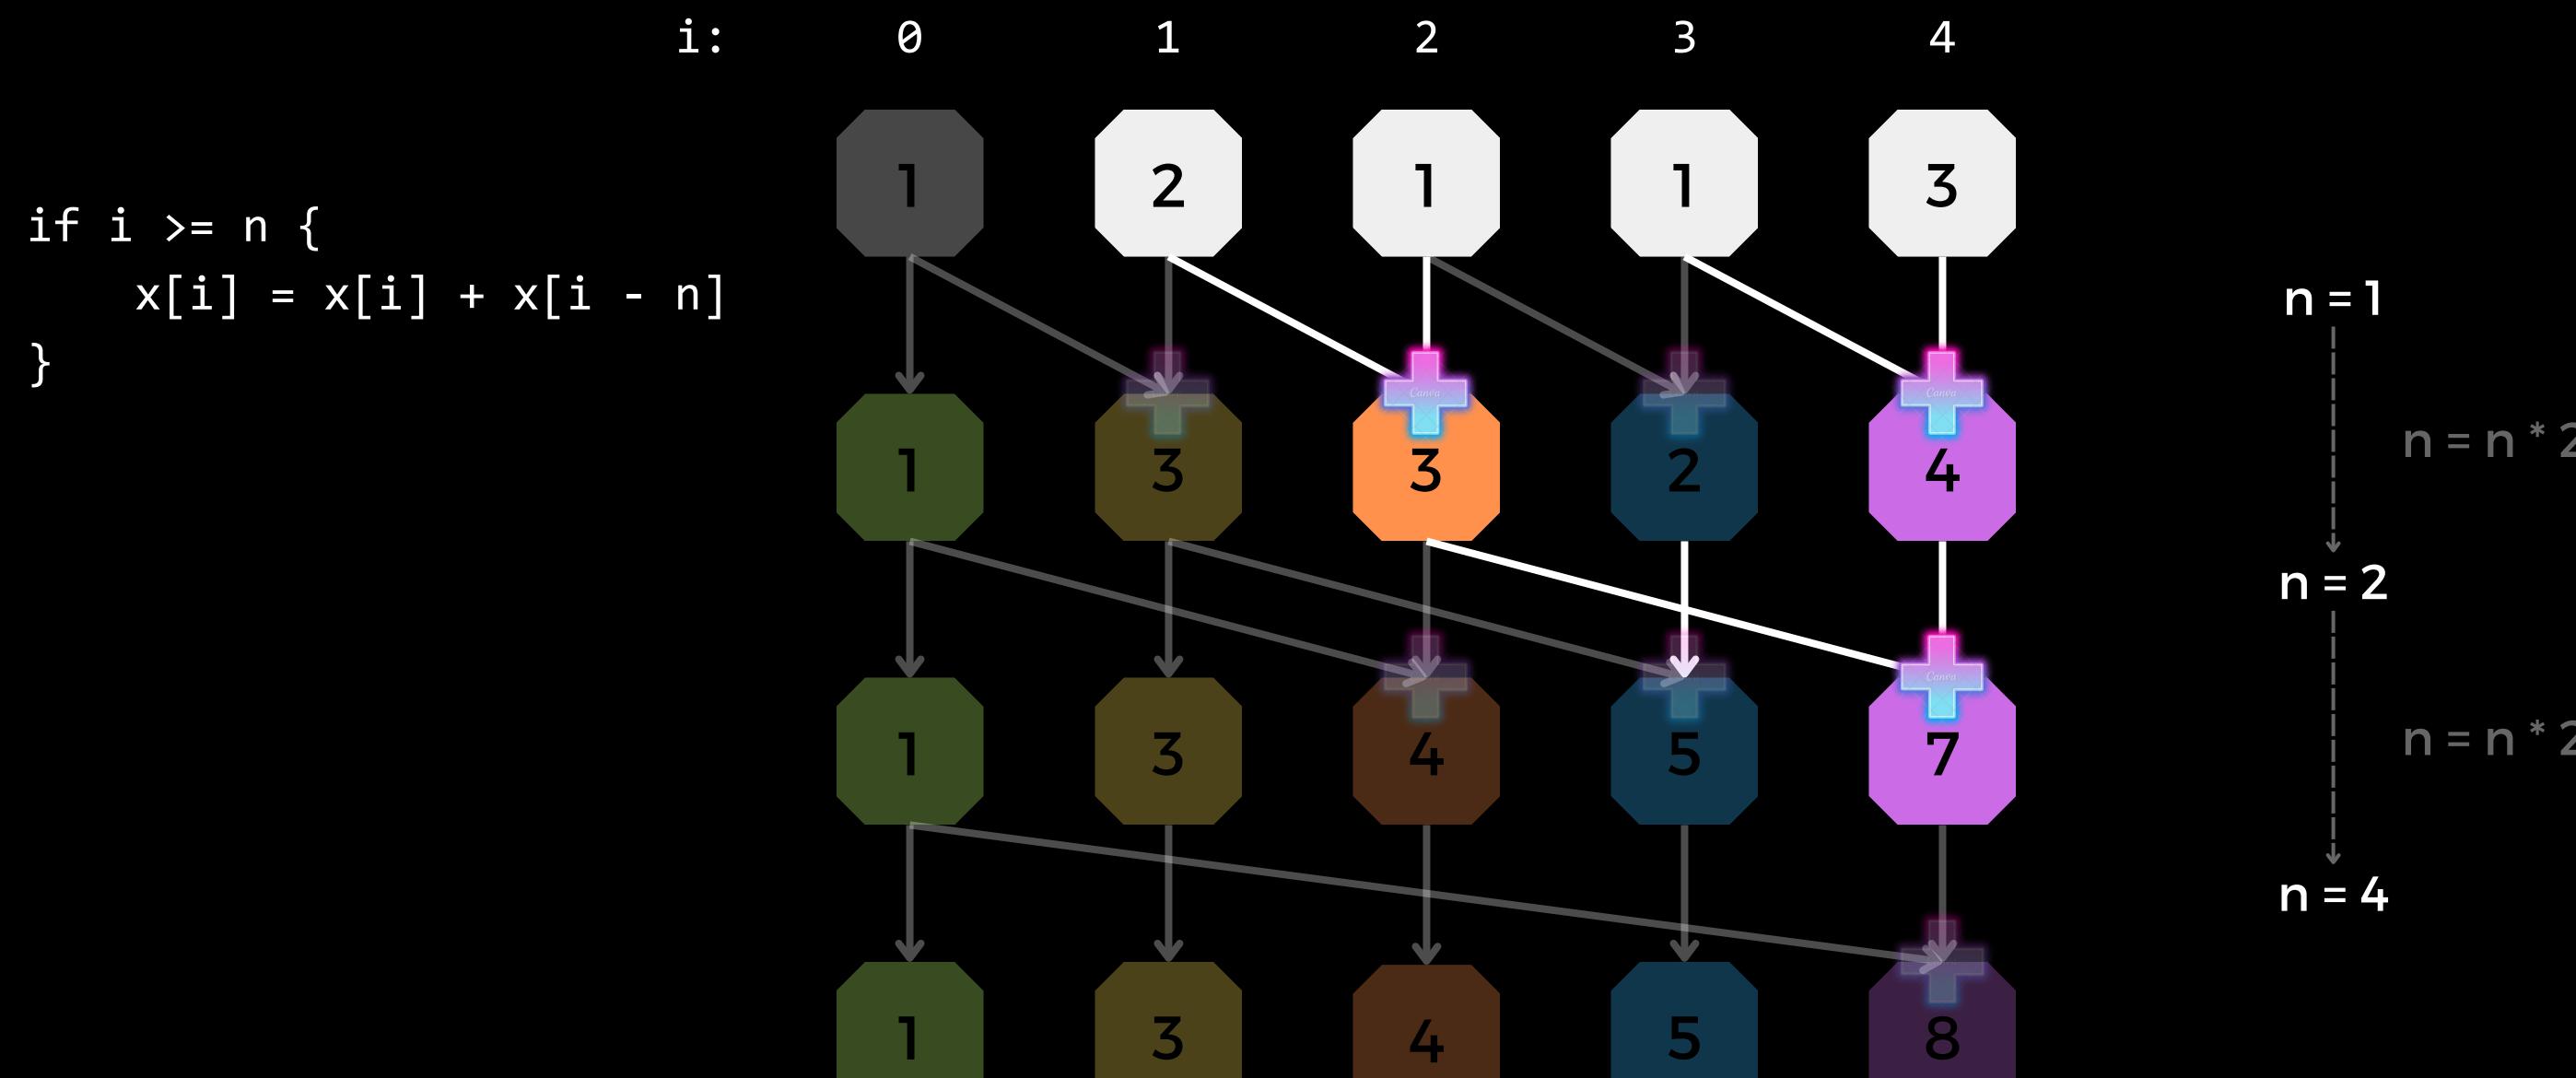
$O(\log_2 n)$

Sum Scan



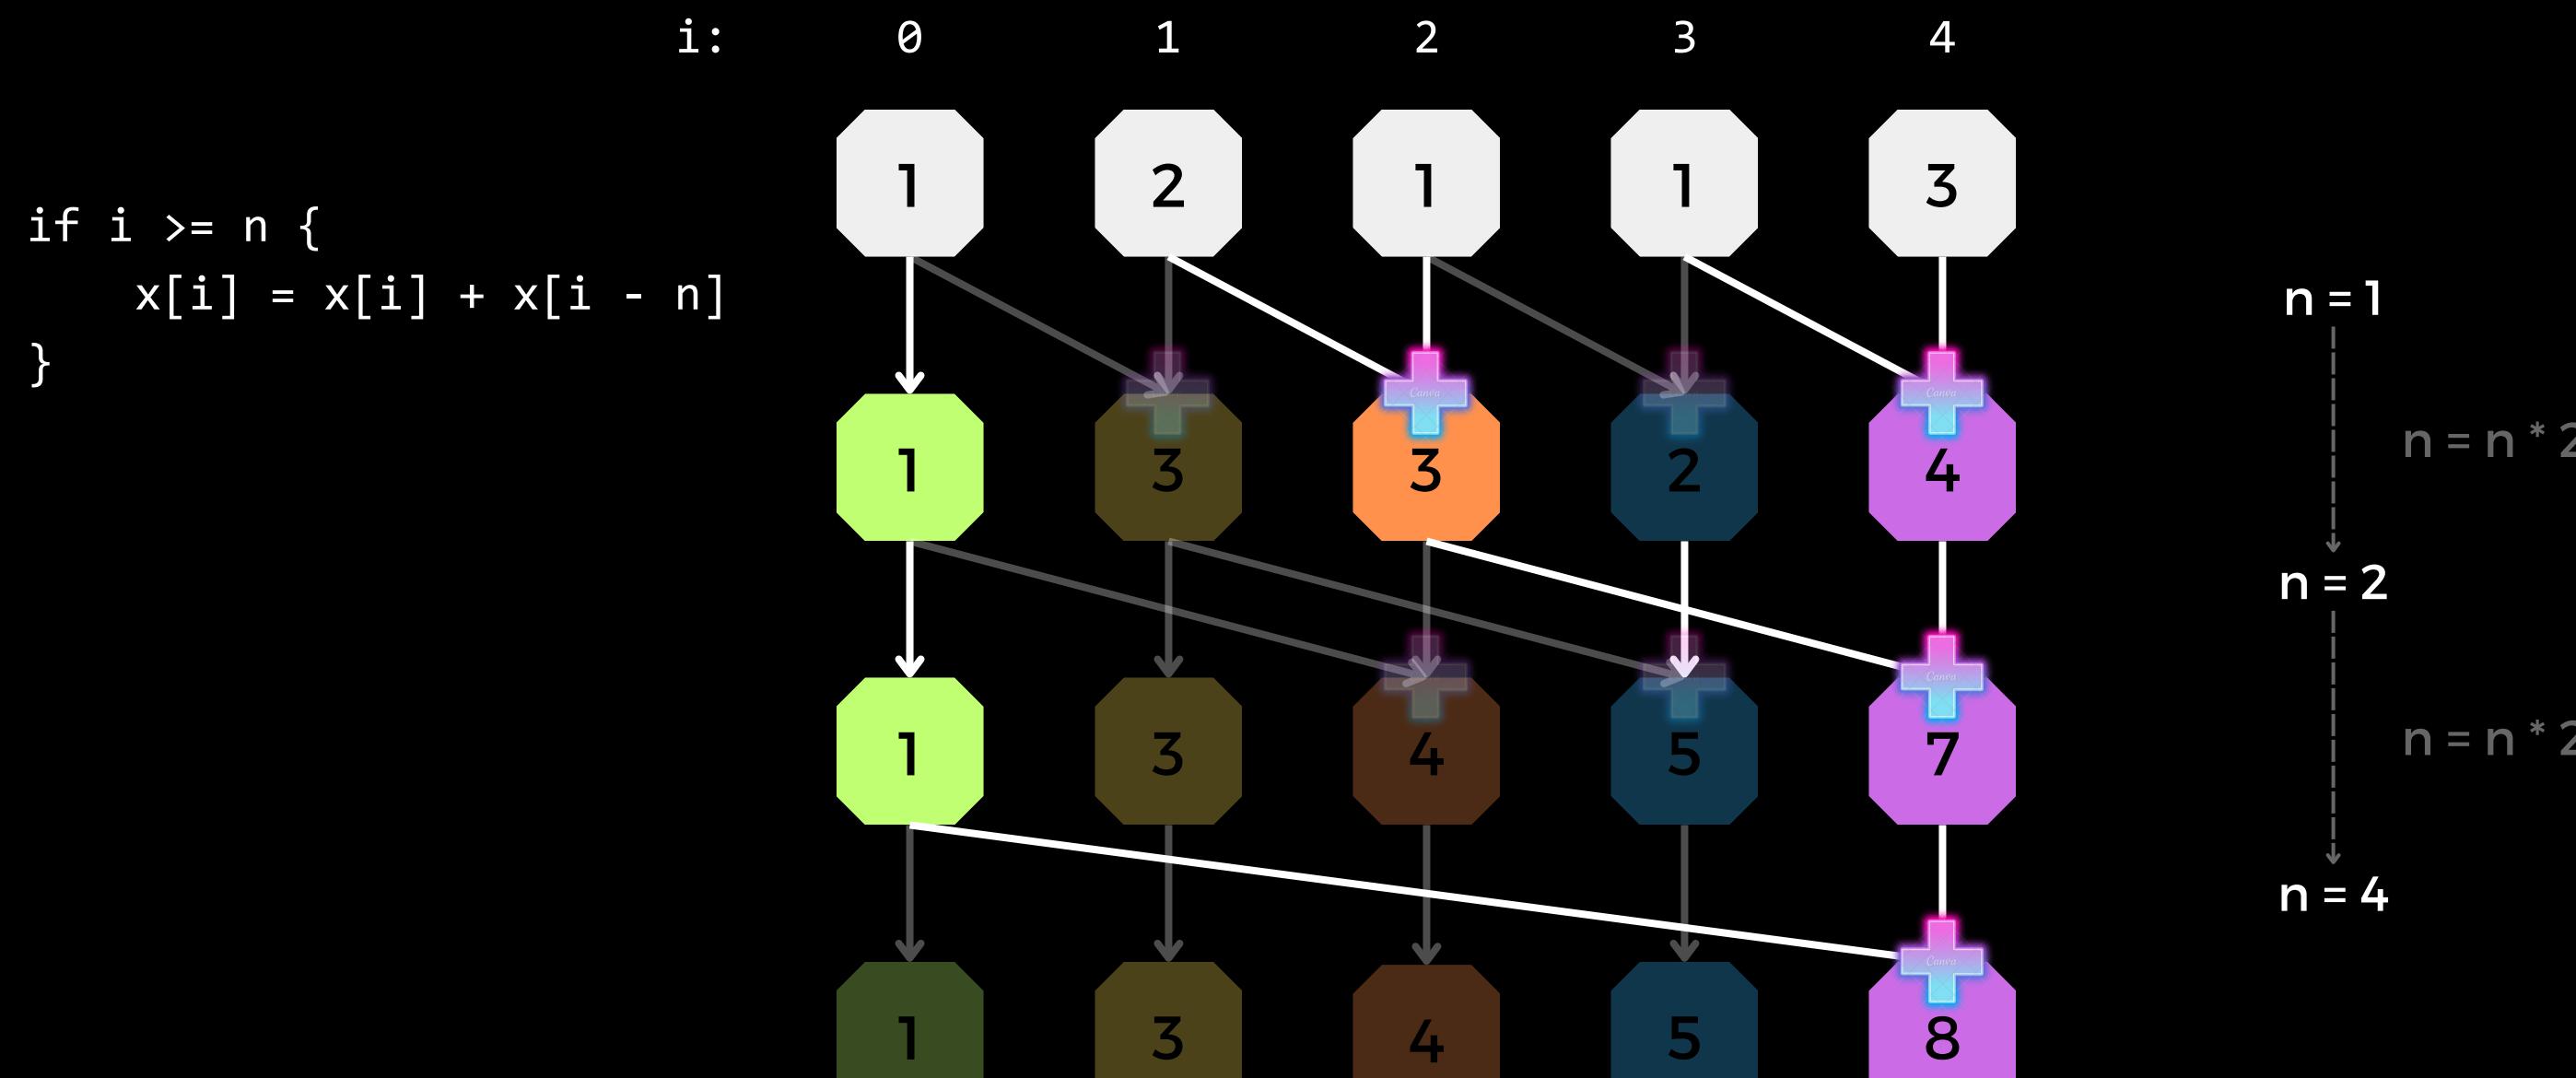
$O(\log_2 n)$

Sum Scan



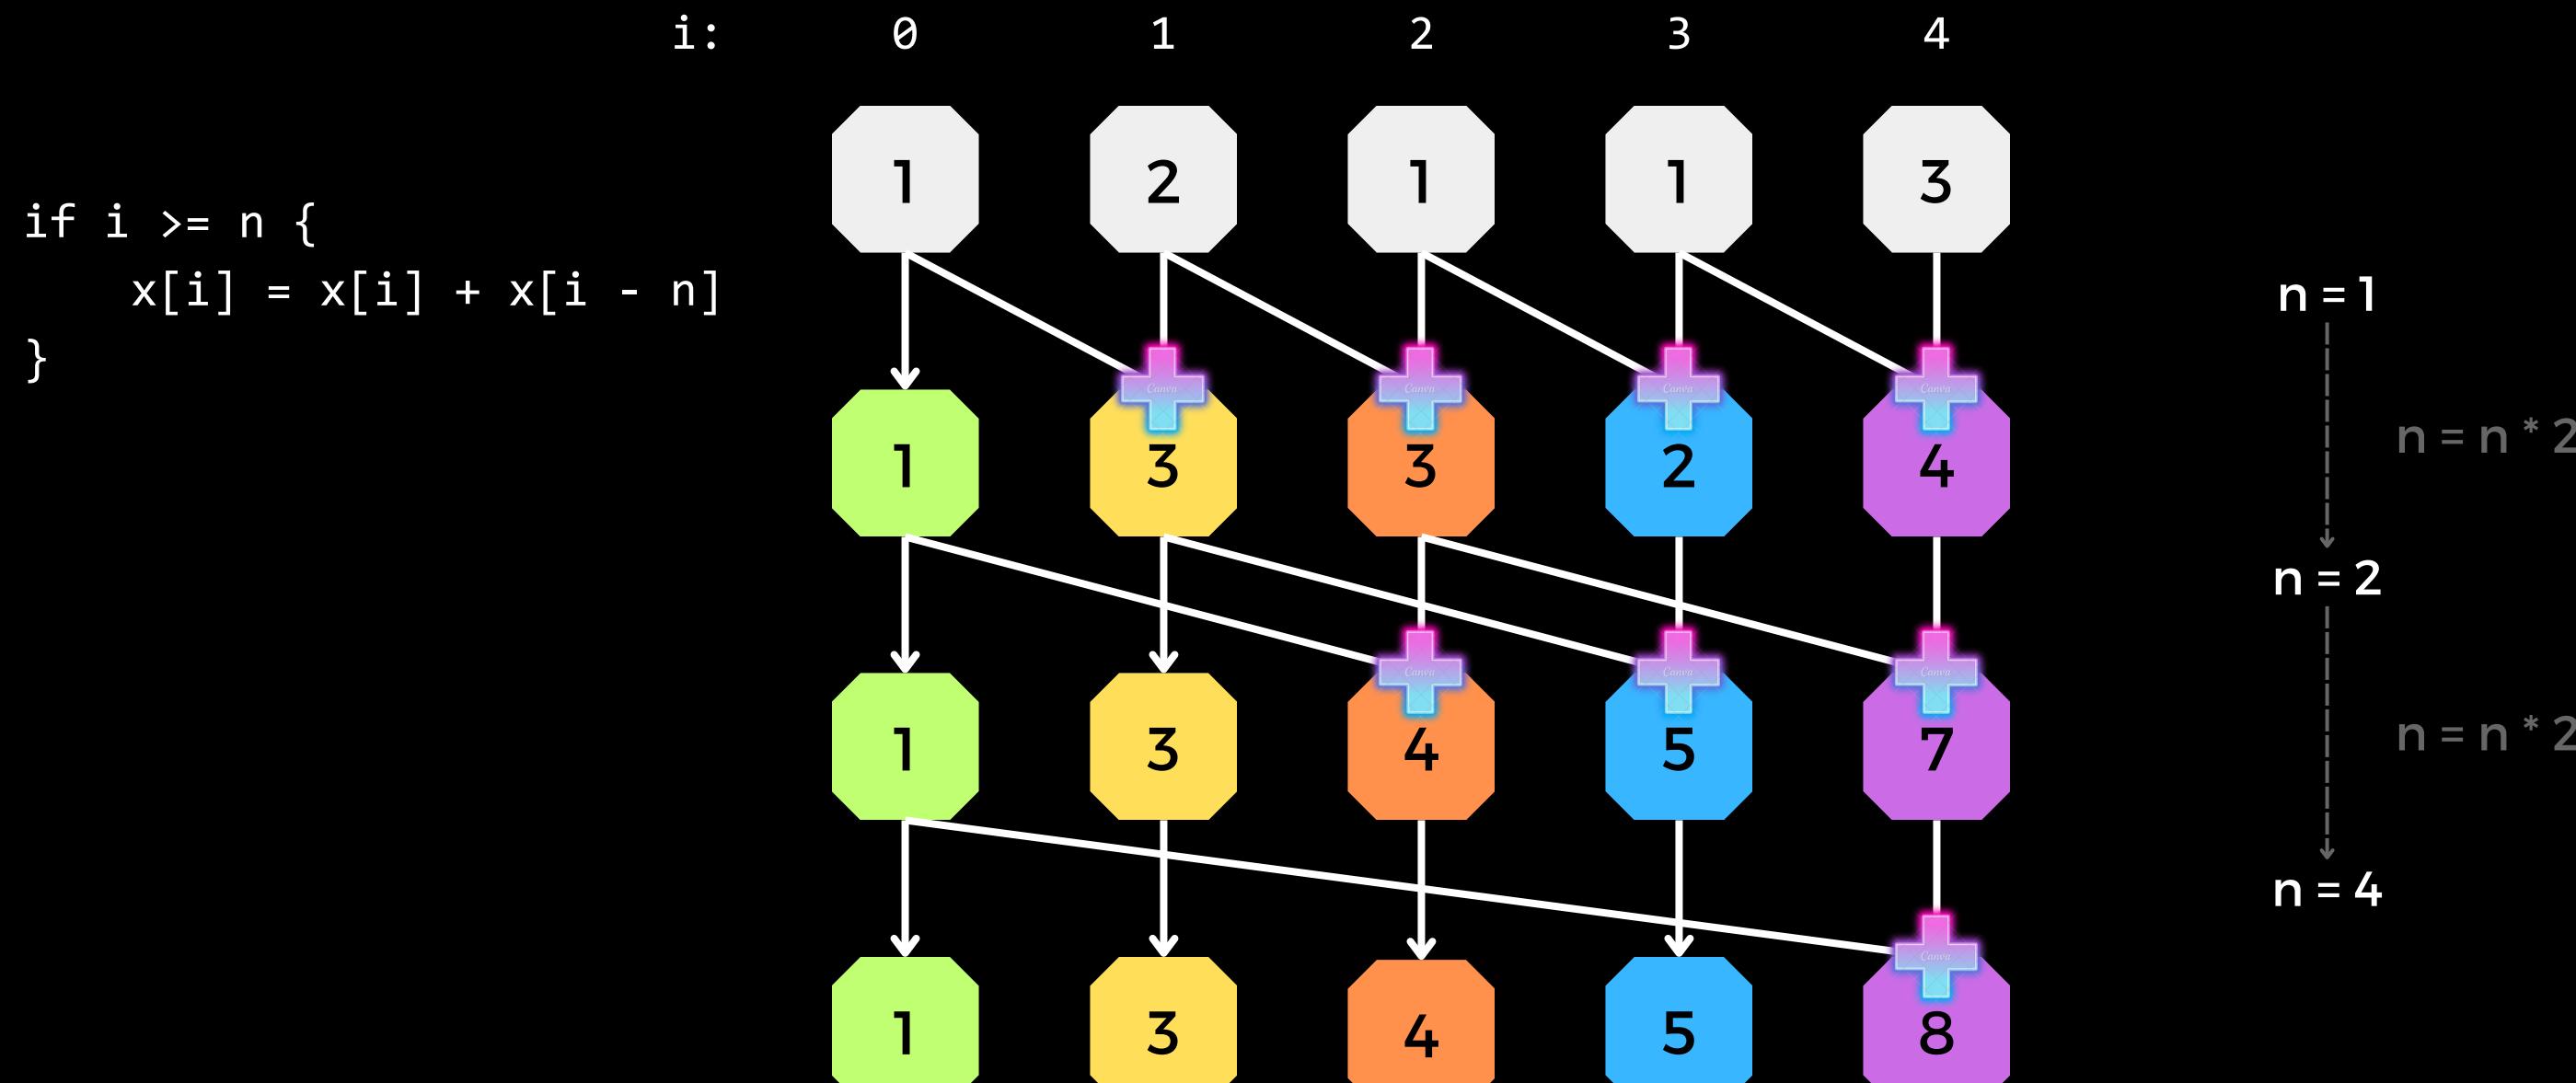
$O(\log_2 n)$

Sum Scan

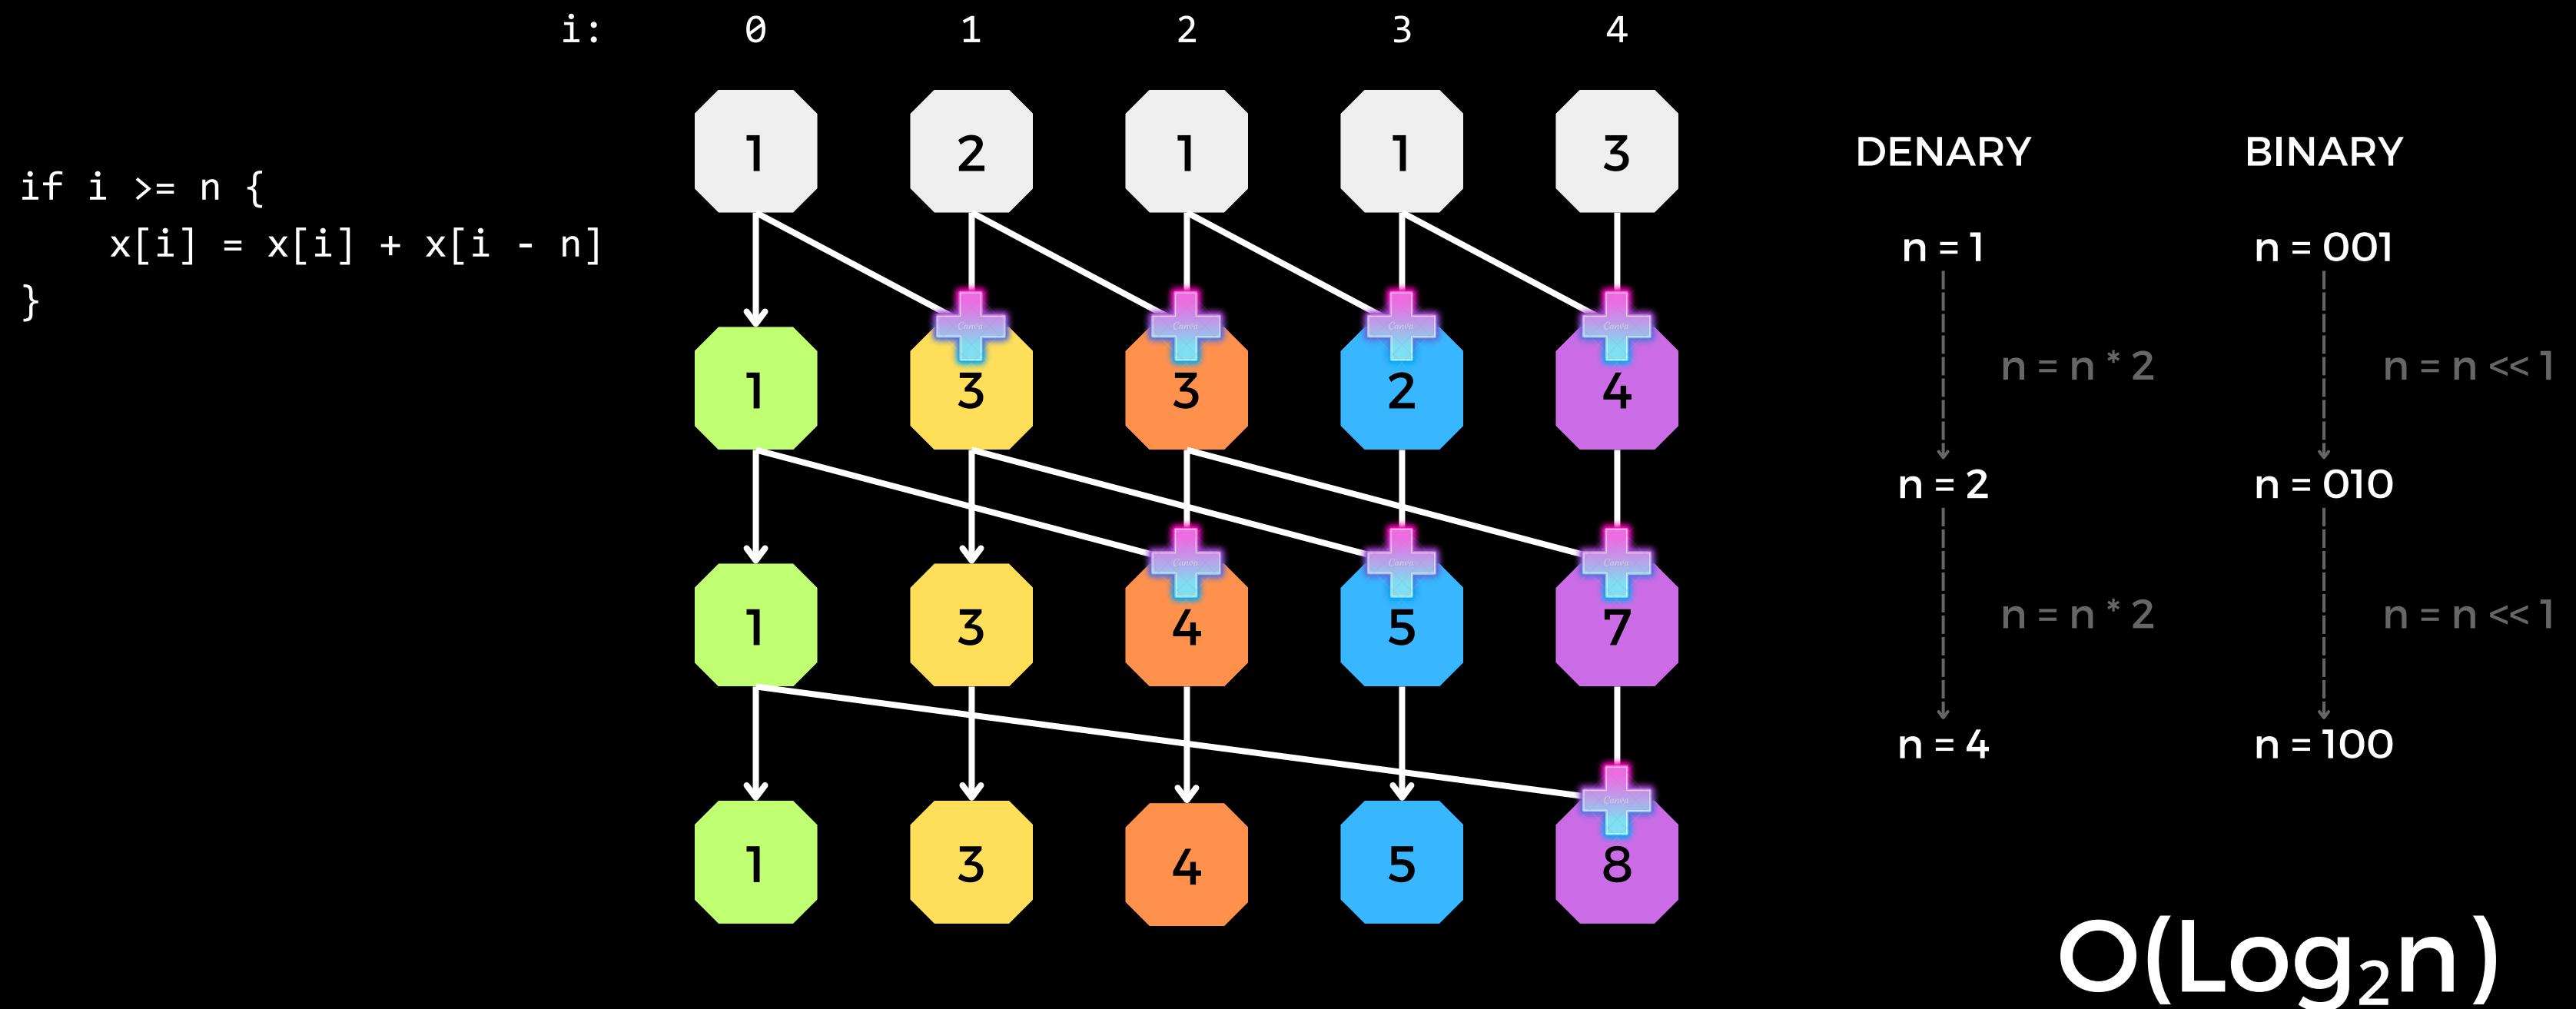


$O(\log_2 n)$

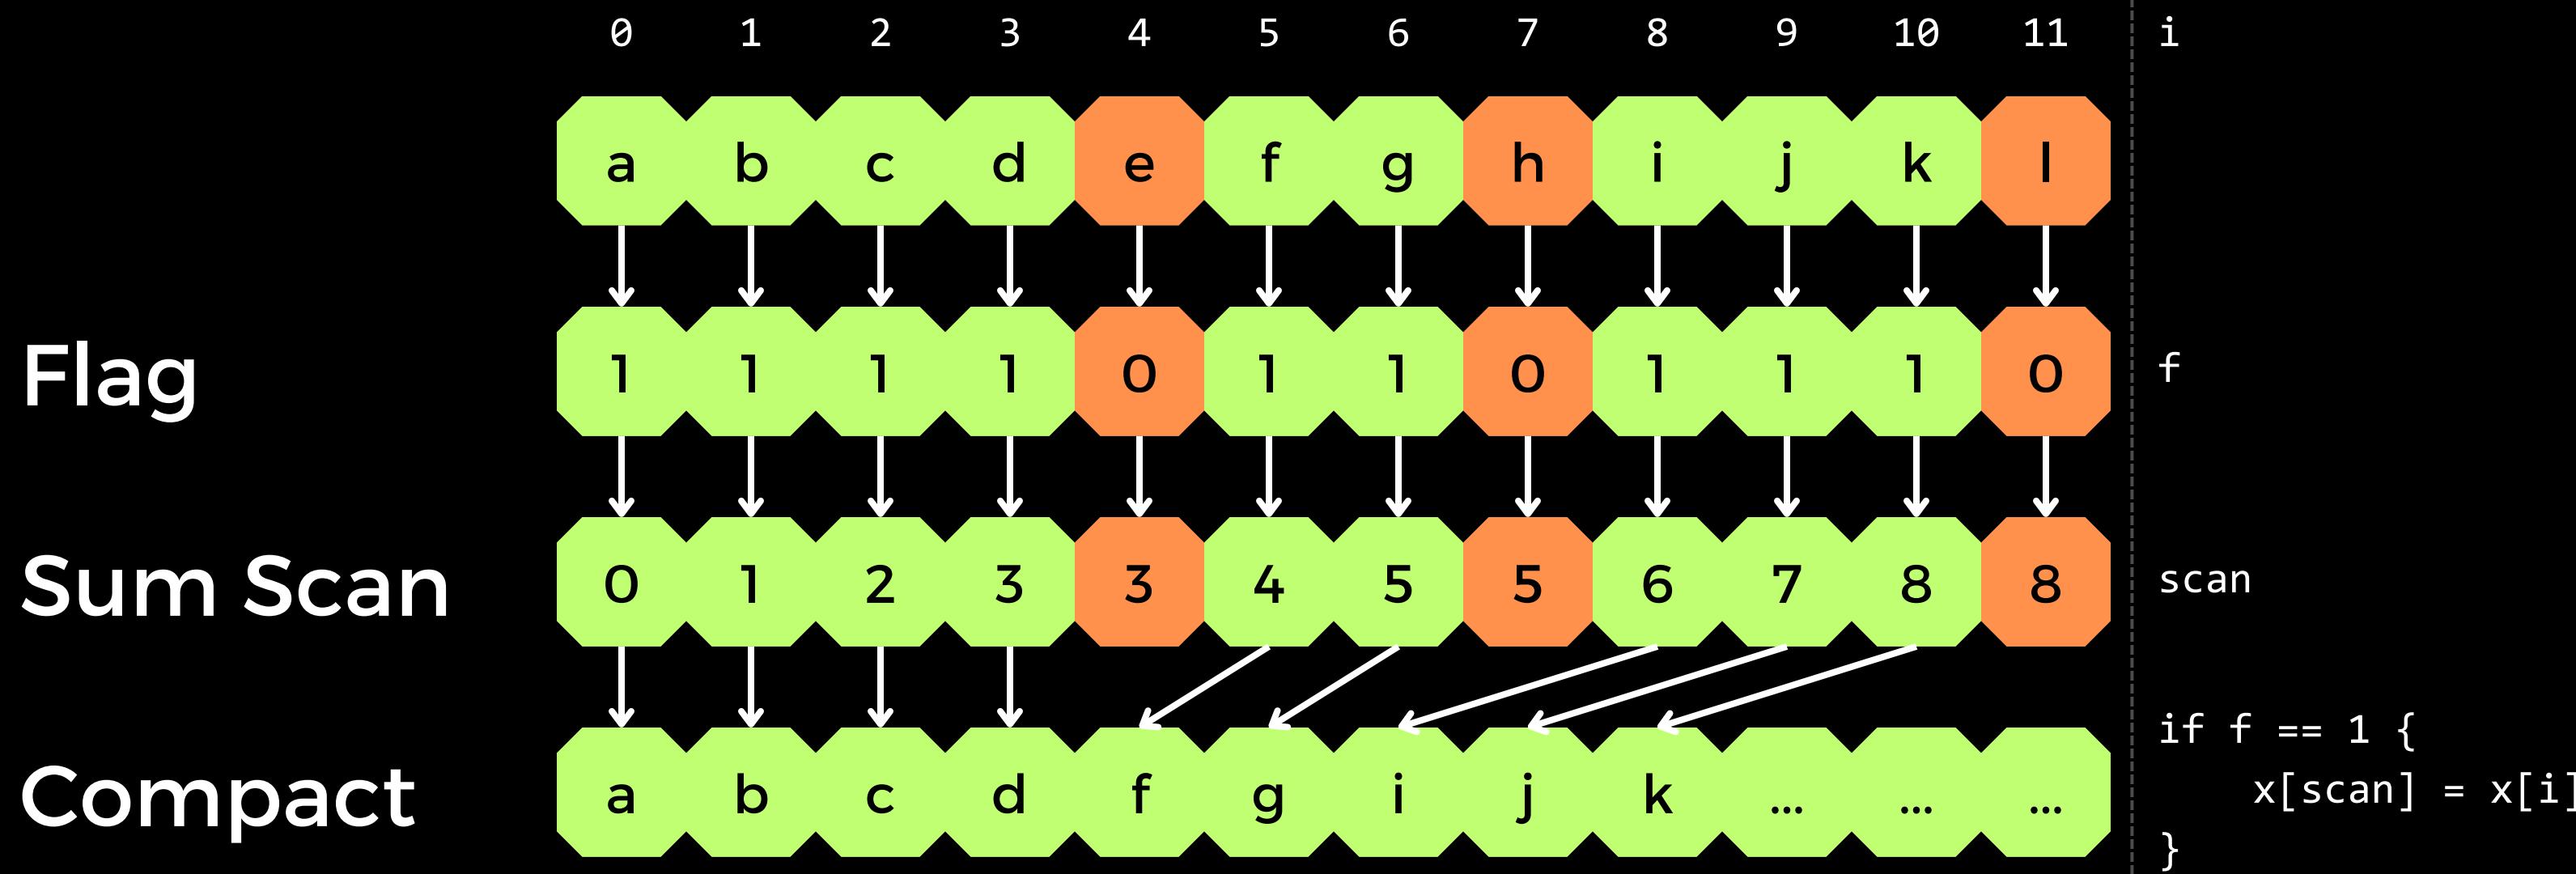
Sum Scan



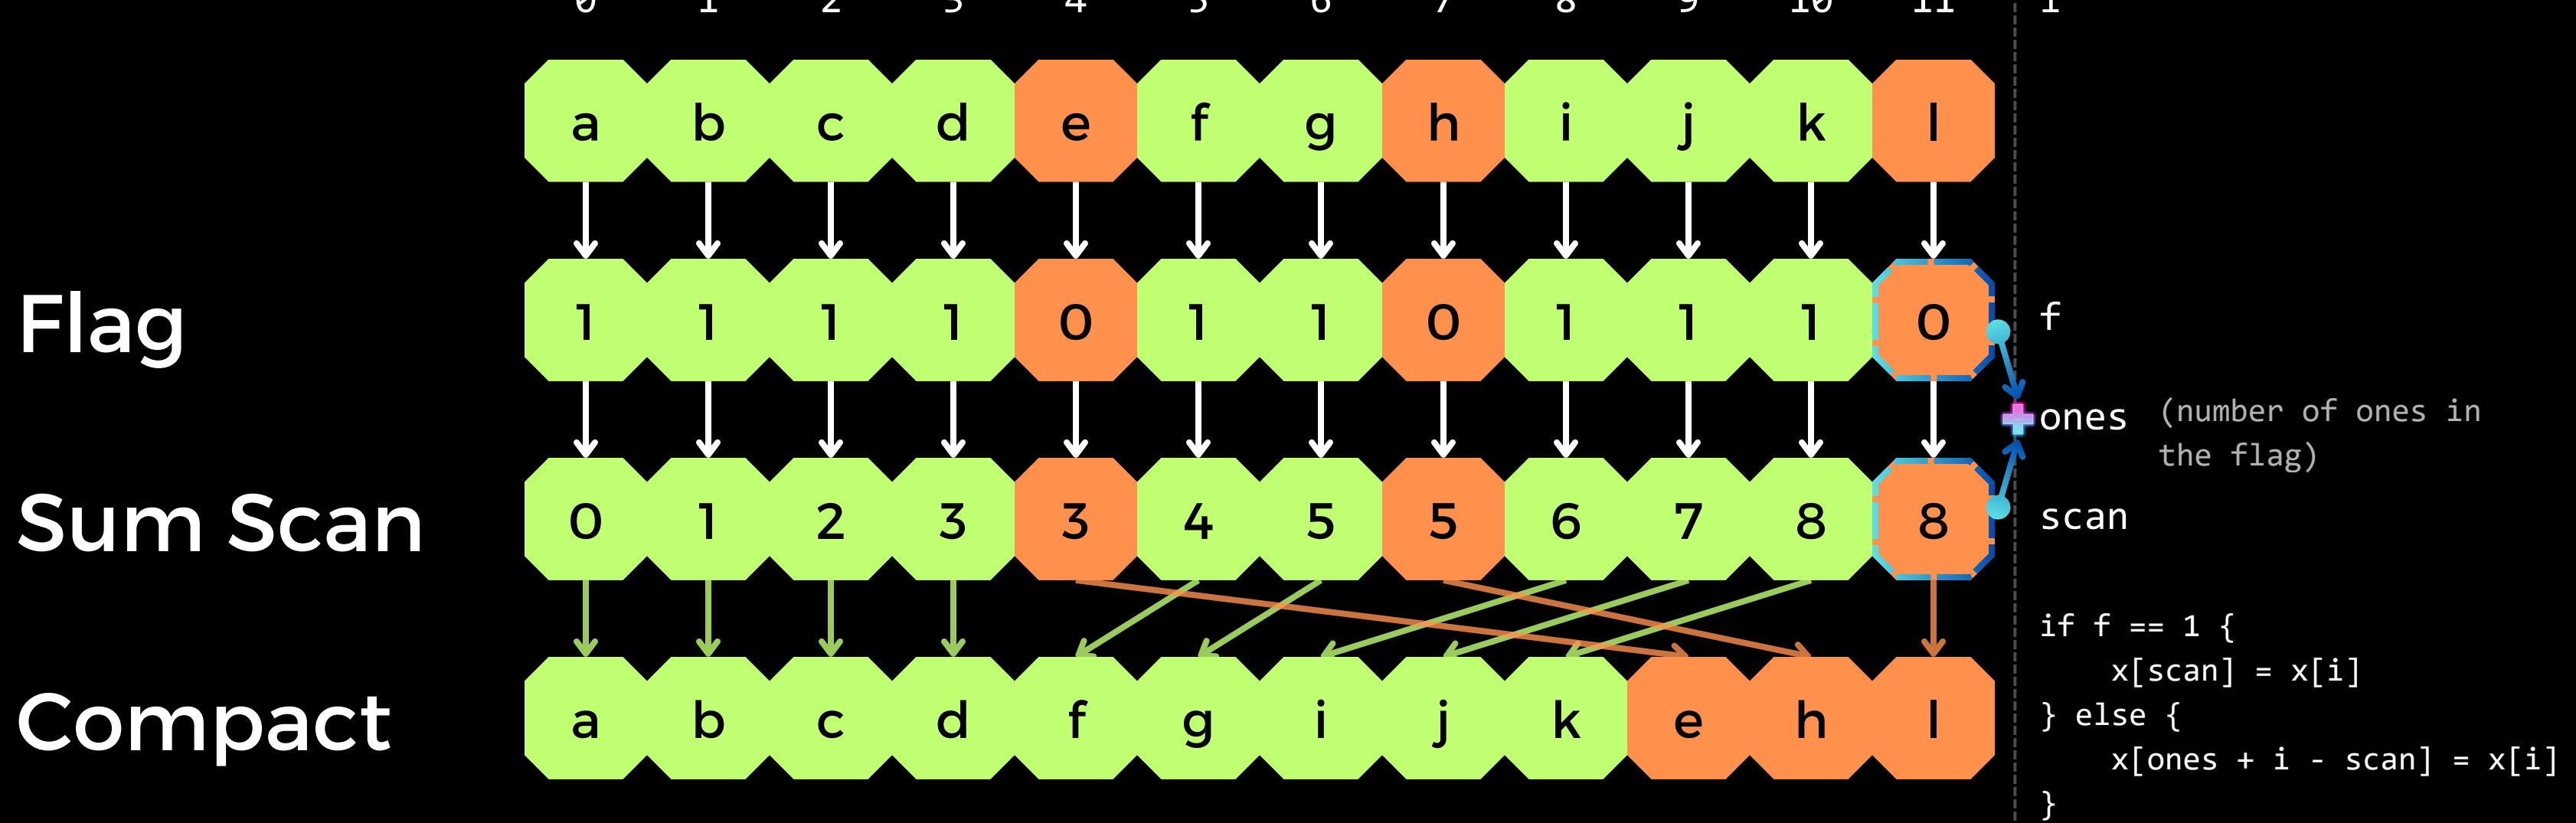
Sum Scan



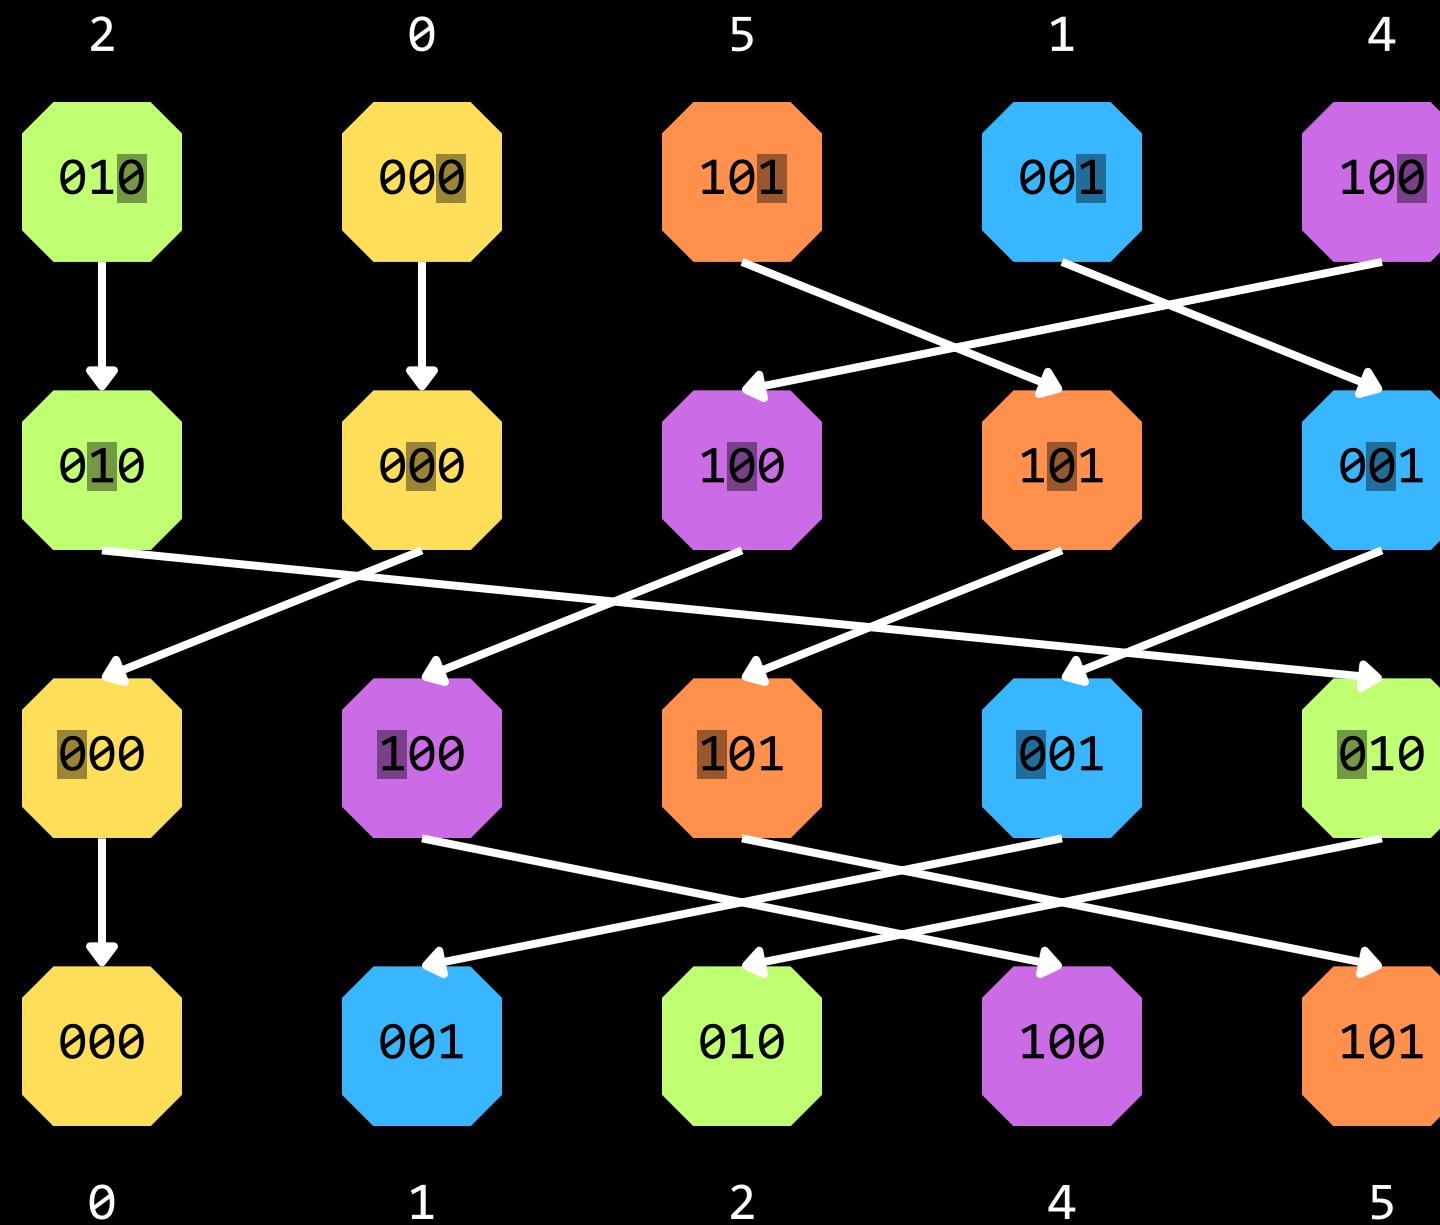
Compaction



Partition

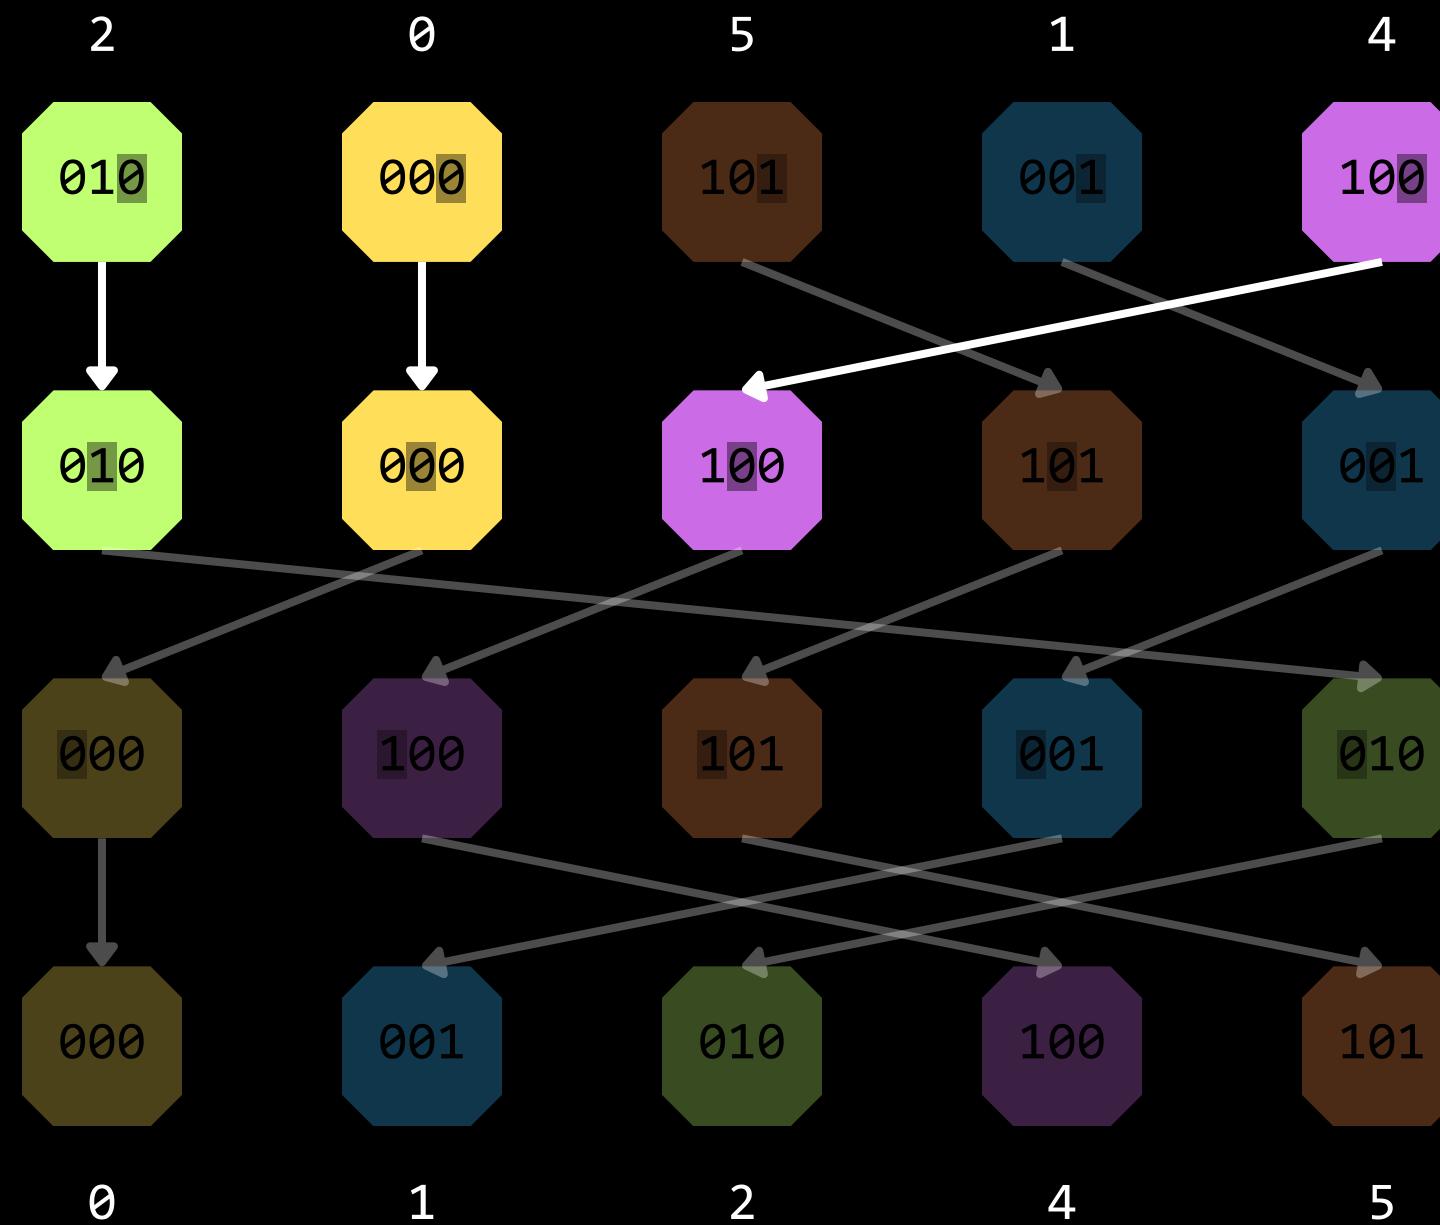


Radix Sort??



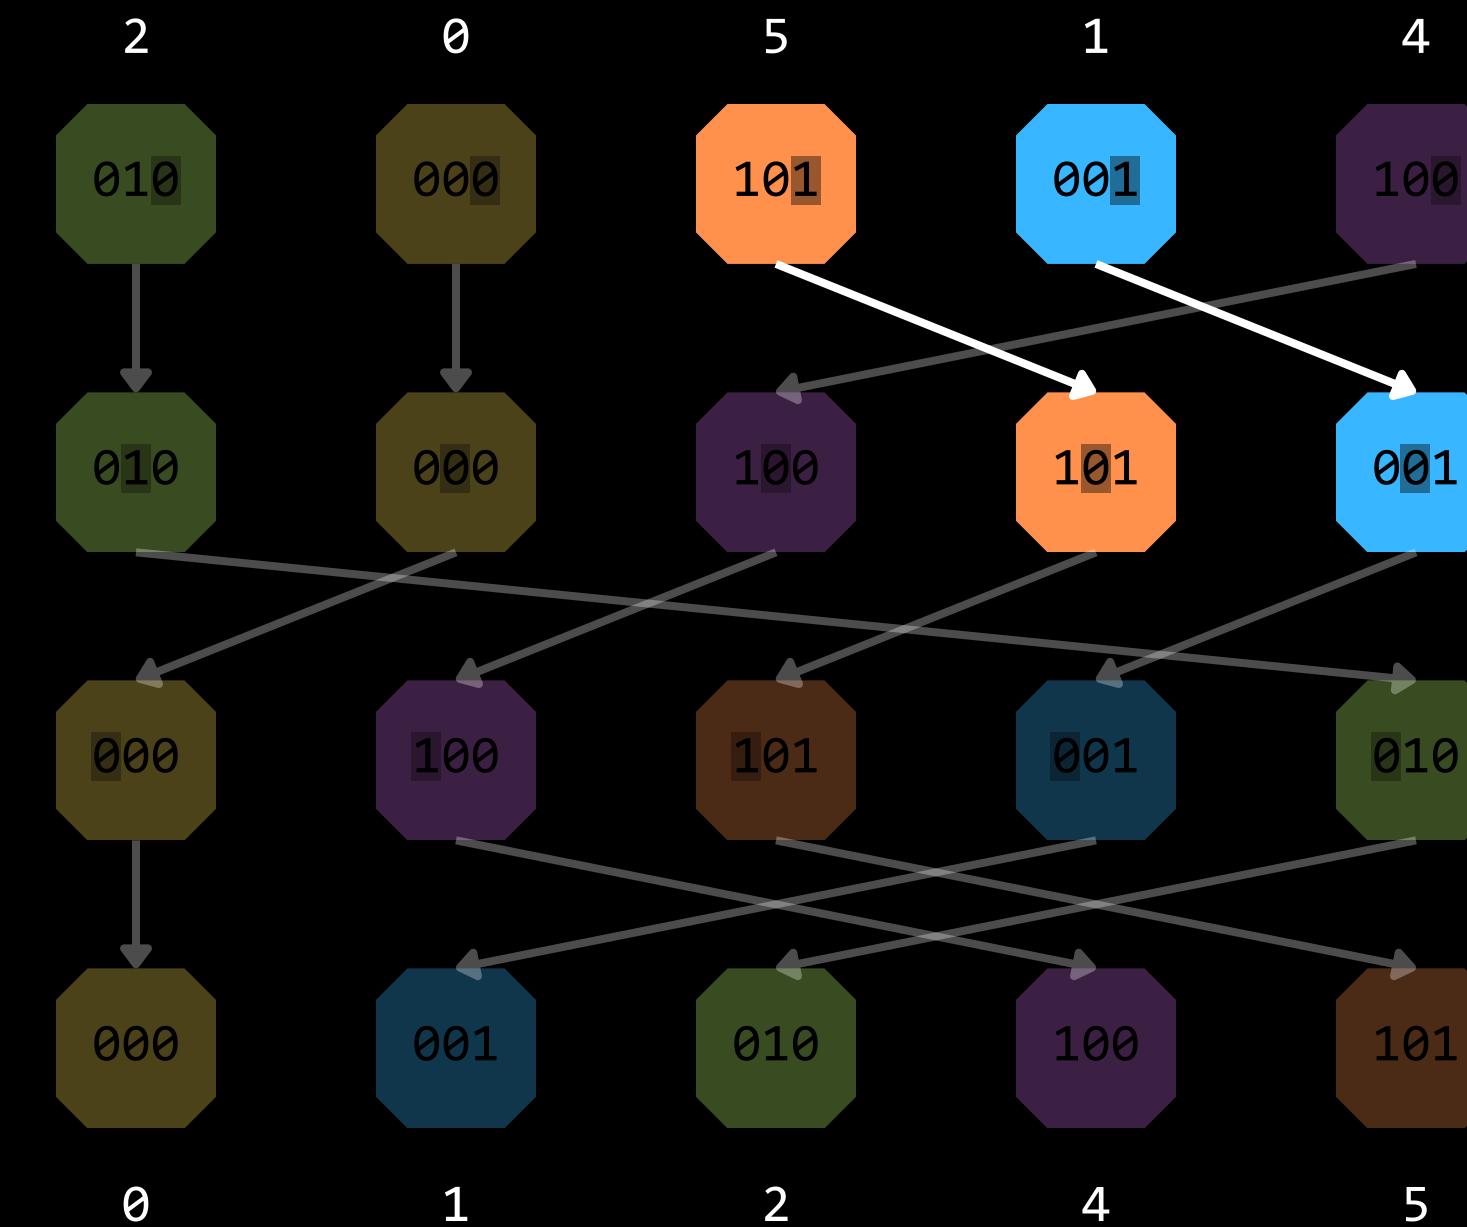
Solve this
using DPPs!

Radix Sort??



Solve this
using DPPs!

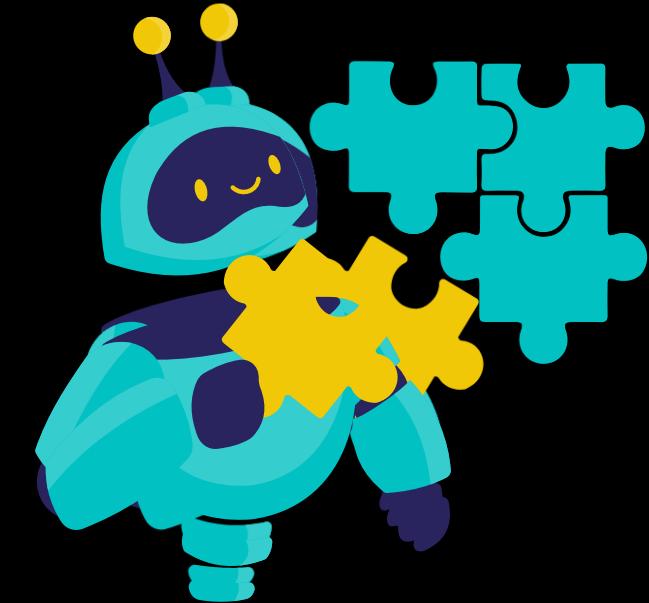
Radix Sort??



Solve this
using DPPs!

CHAPTER 4

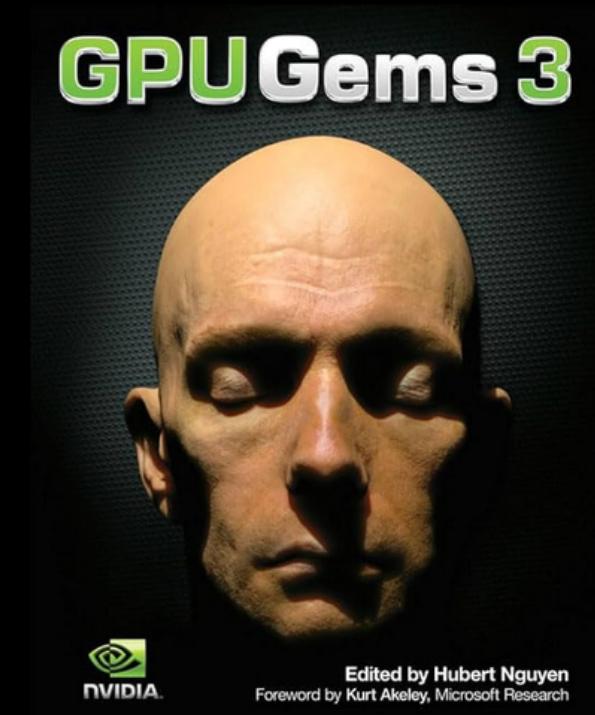
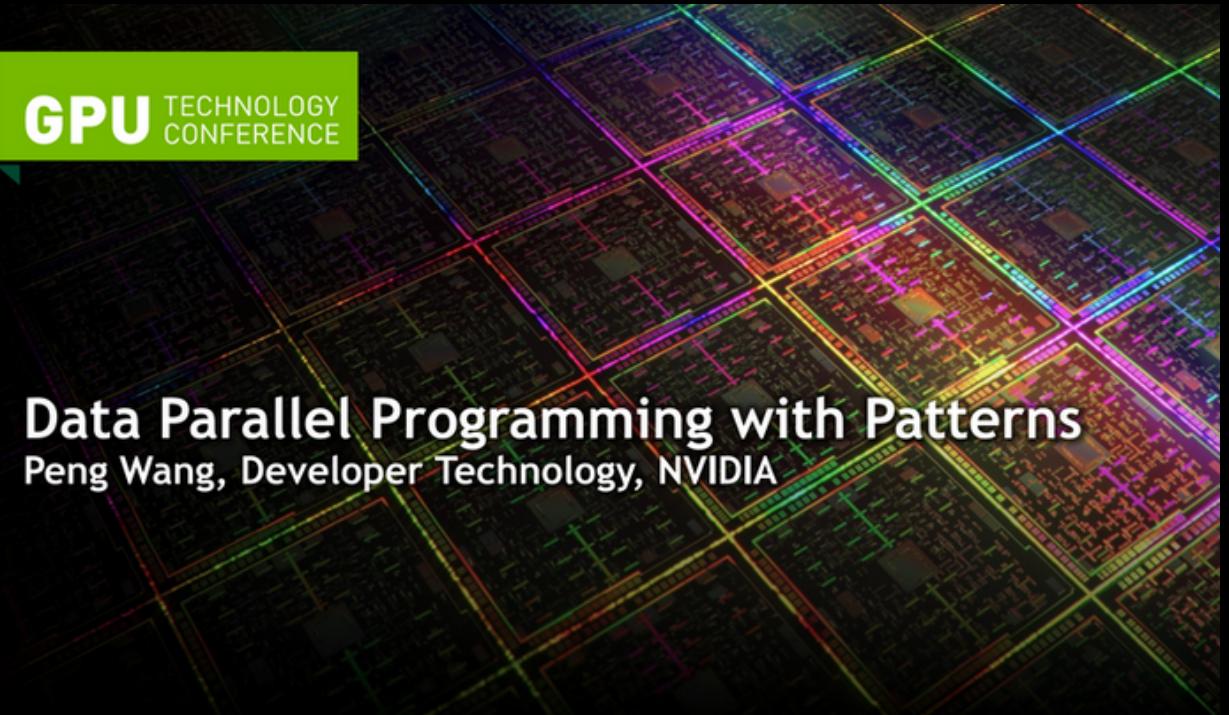
CONCLUSION



Look out!

- Thread count (hardware is still important)
 - Write/mutate data (utilize atomics if necessary)
 - Graph coloring
 - Memory is slow
 - Beware of branching
 - Data structures
 - Use DPP to express your algorithm
-

Additional Resources



Q & A?



VOXELL



THANK YOU

HAVE FUN PARALLELIZING!



VOXELL



